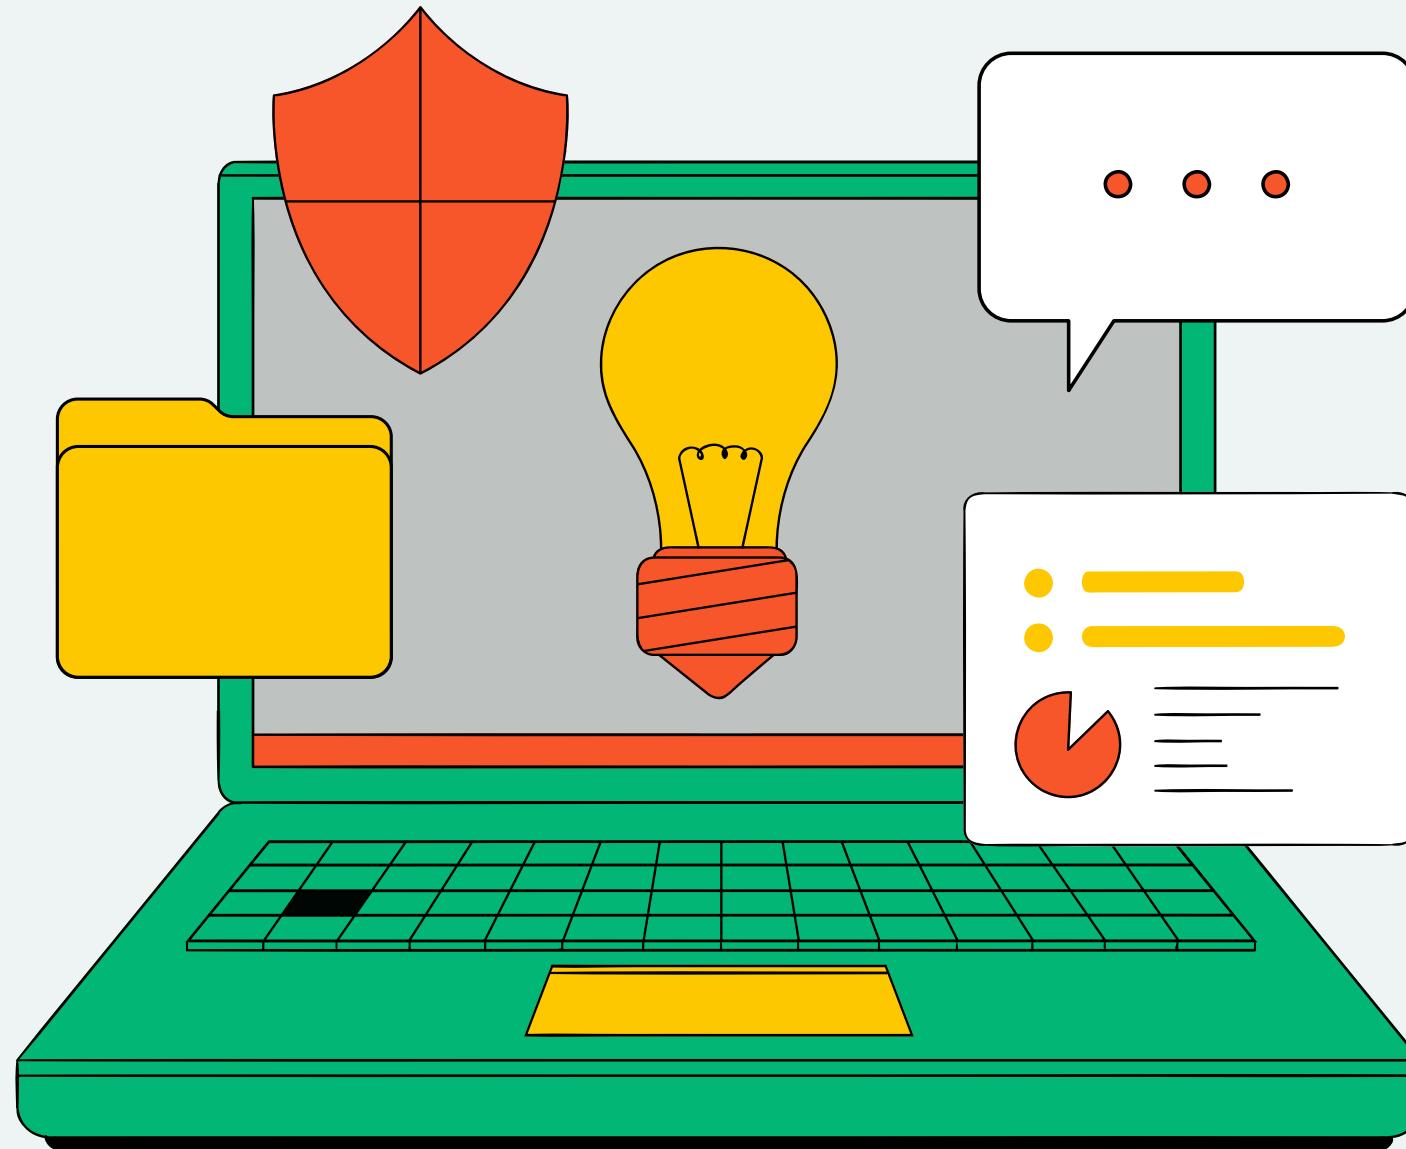
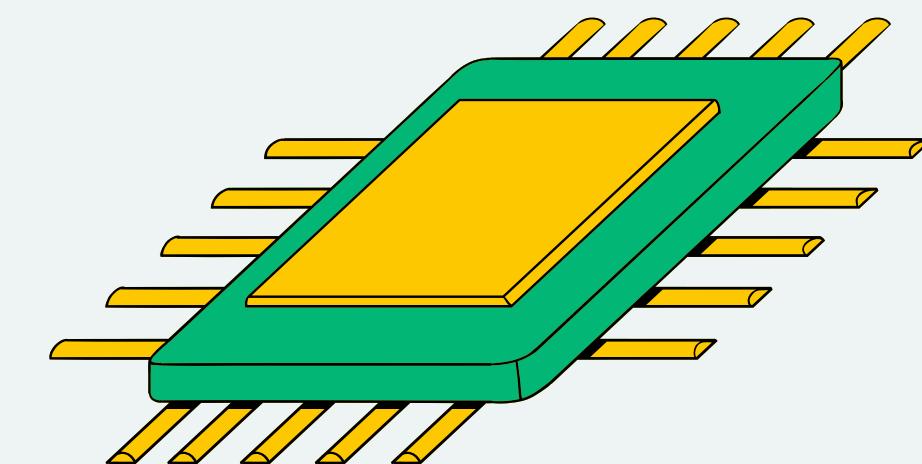


CS 251, MACHINE LEARNING

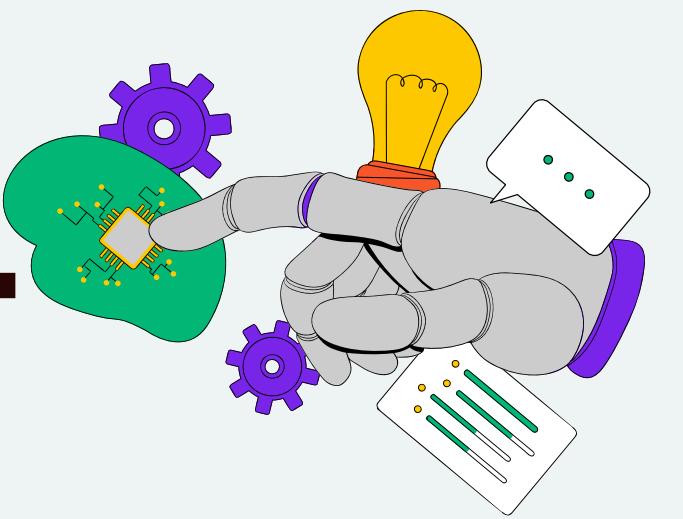


# IMAGE CLASSIFICATION

Group Members:  
Narek Ghukasyan, Shushan Gevorgyan, Armen Madoyan, Yeva Stepanyan



# DATASET



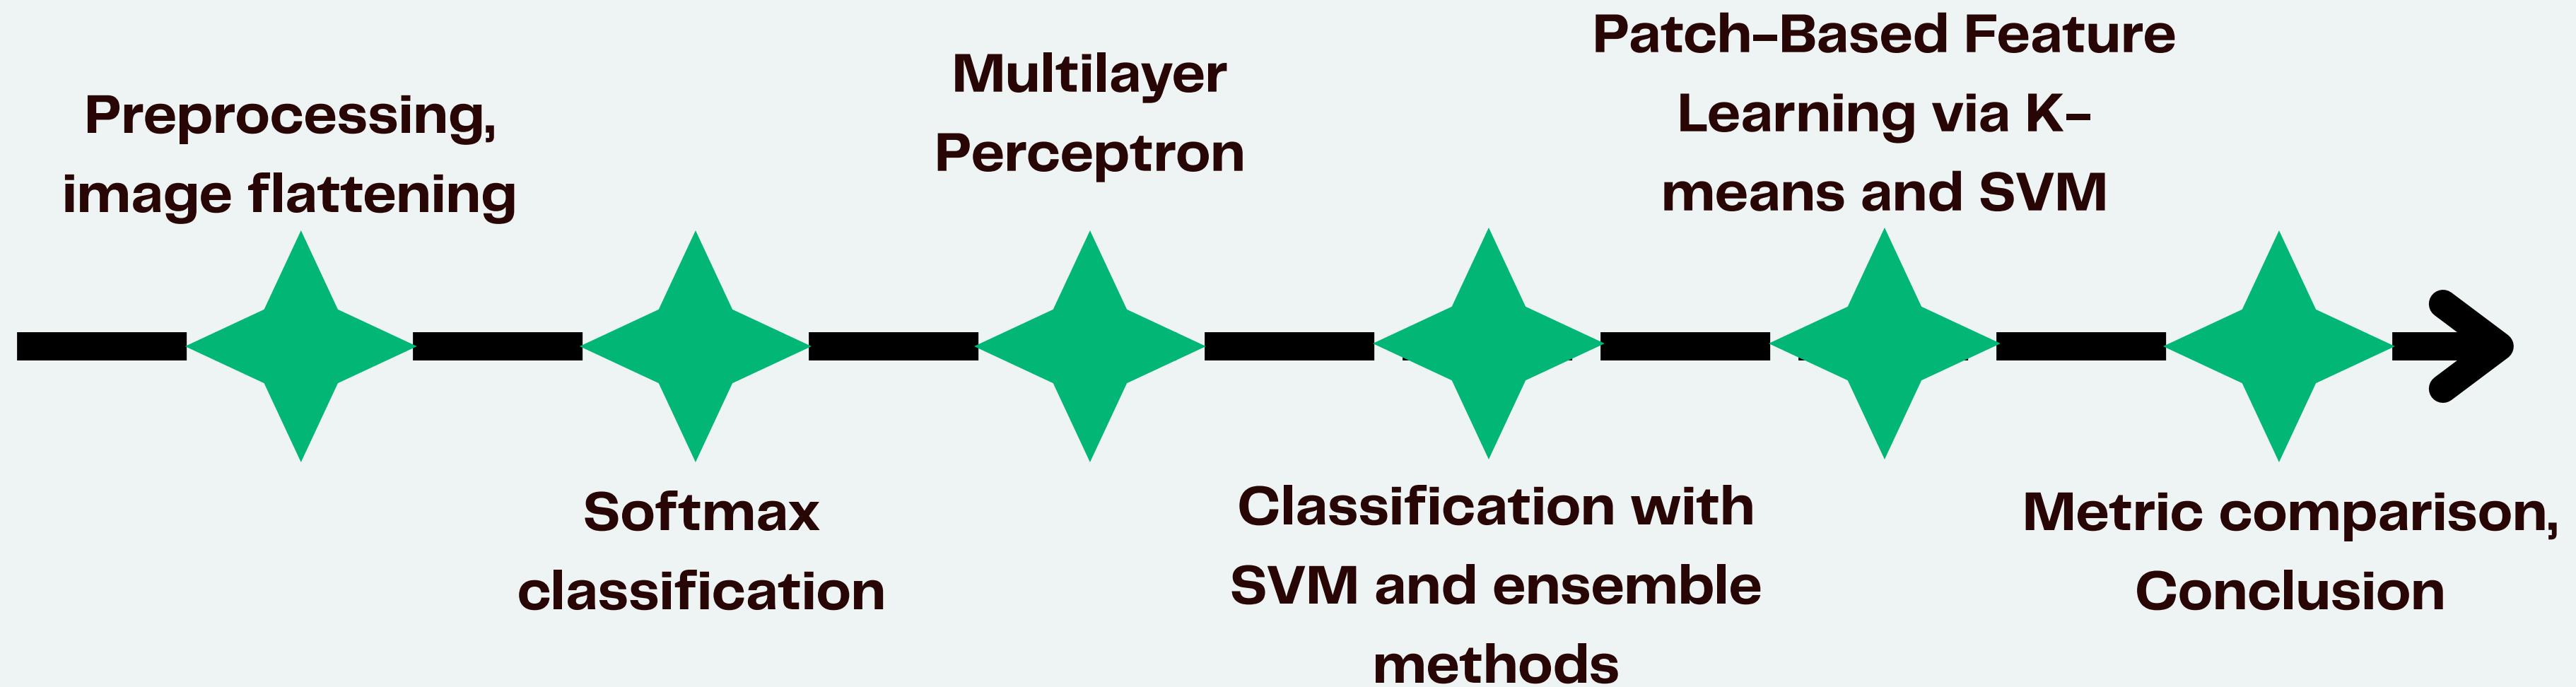
|            |  |  |  |  |  |  |  |  |  |  |  |
|------------|--|--|--|--|--|--|--|--|--|--|--|
| airplane   |  |  |  |  |  |  |  |  |  |  |  |
| automobile |  |  |  |  |  |  |  |  |  |  |  |
| bird       |  |  |  |  |  |  |  |  |  |  |  |
| cat        |  |  |  |  |  |  |  |  |  |  |  |
| deer       |  |  |  |  |  |  |  |  |  |  |  |
| dog        |  |  |  |  |  |  |  |  |  |  |  |
| frog       |  |  |  |  |  |  |  |  |  |  |  |
| horse      |  |  |  |  |  |  |  |  |  |  |  |
| ship       |  |  |  |  |  |  |  |  |  |  |  |
| truck      |  |  |  |  |  |  |  |  |  |  |  |

## CIFAR-10

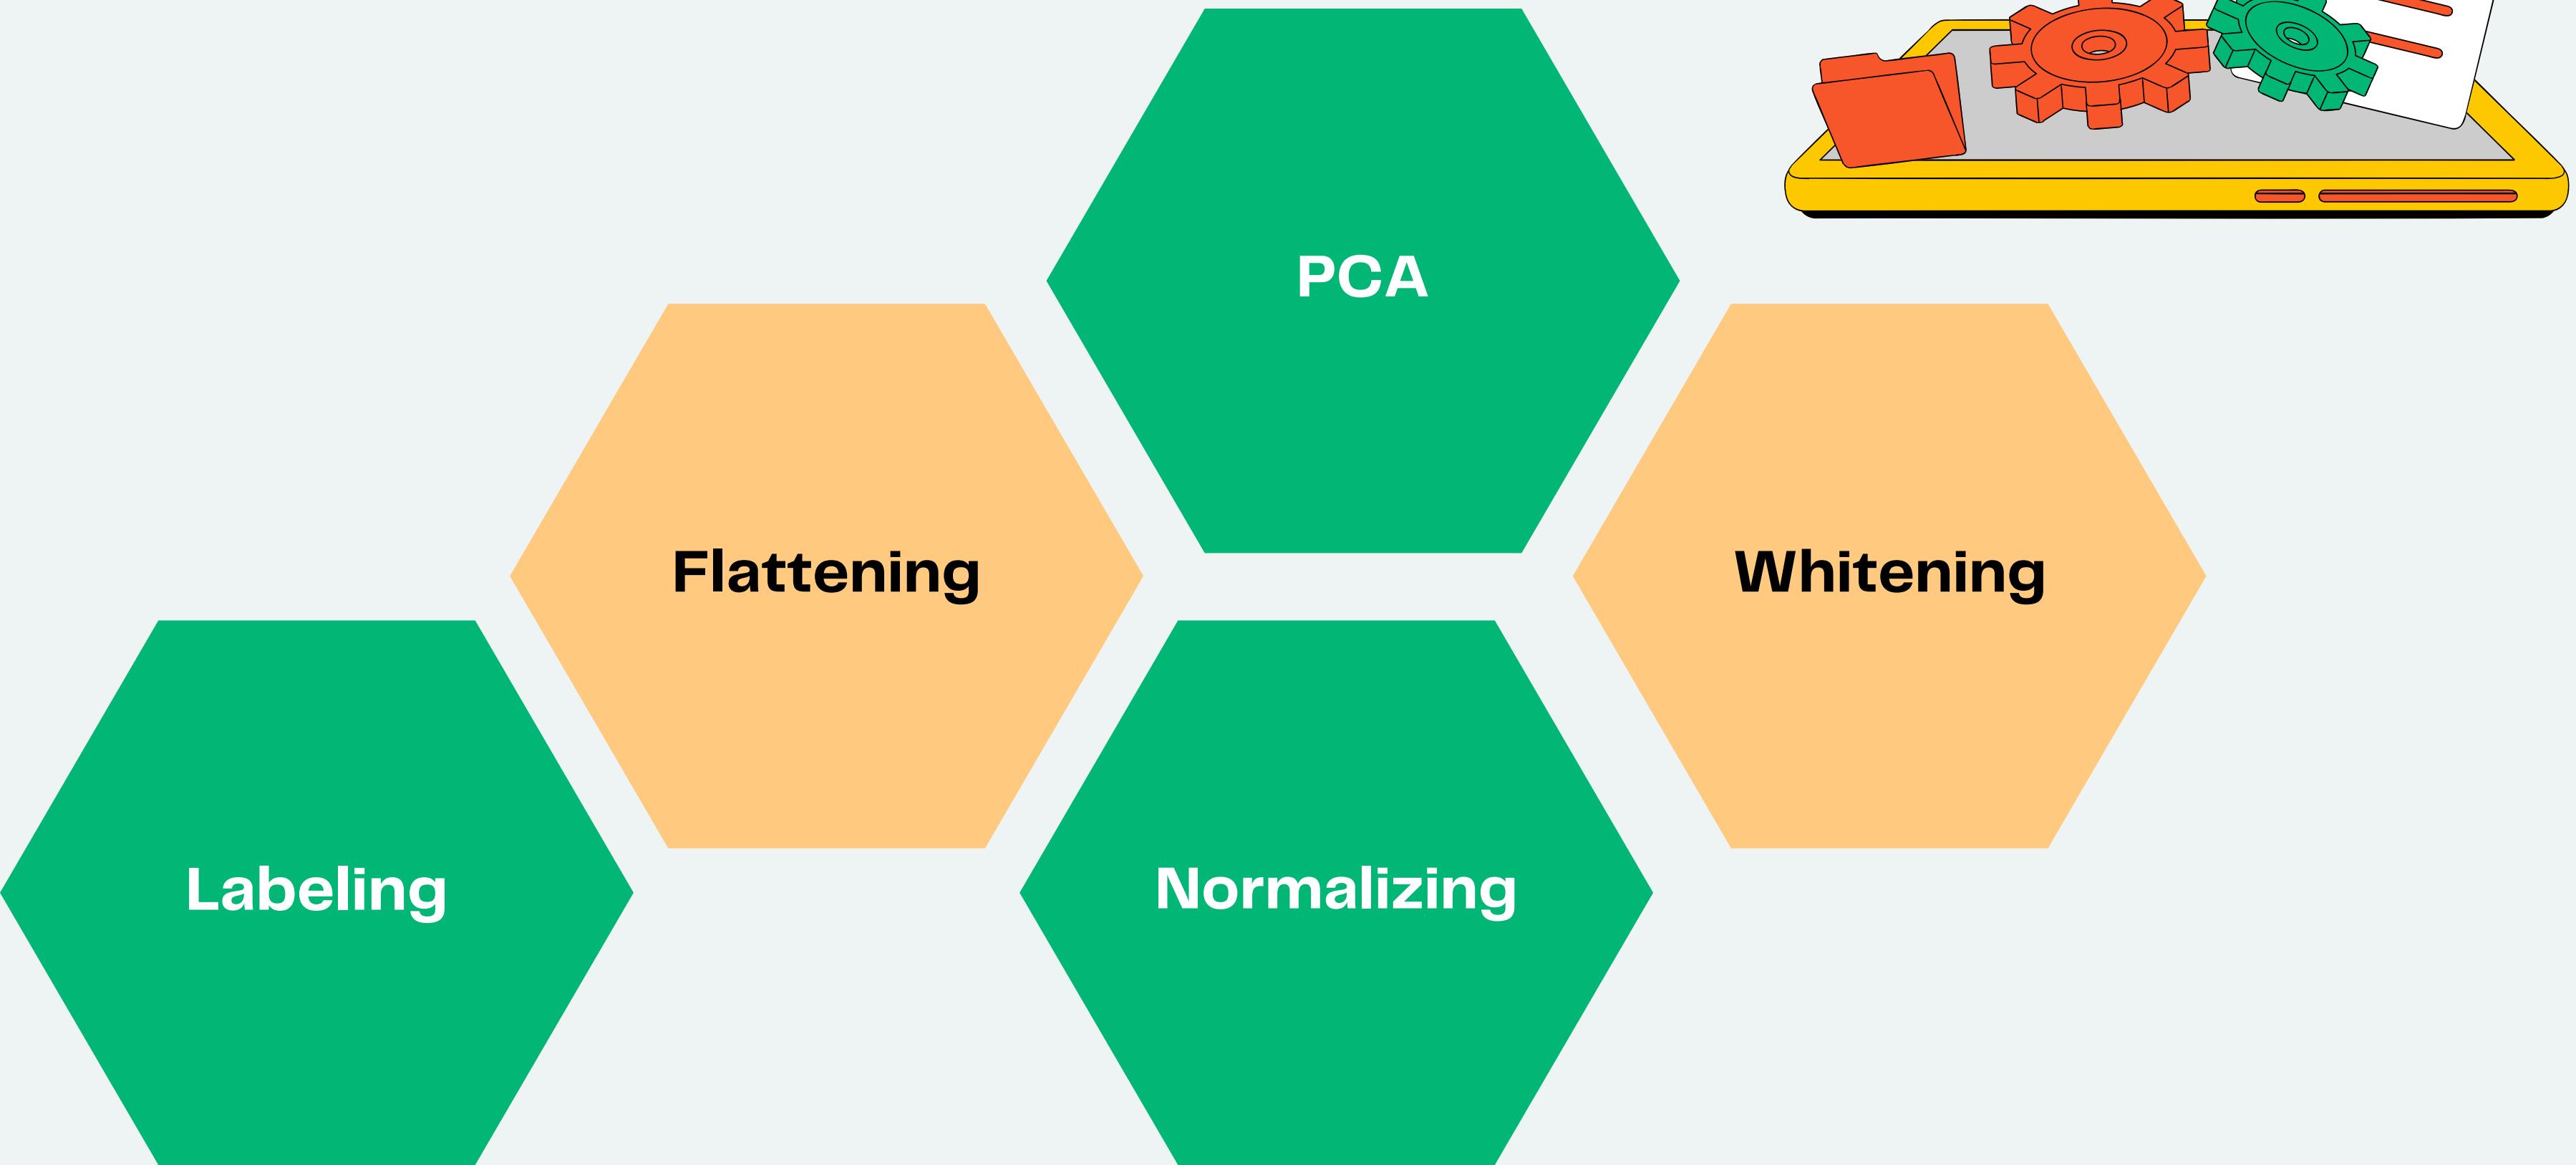
32x32 colour images in 10 classes

60.000 images, 6000 image per class  
(the dataset is balanced)

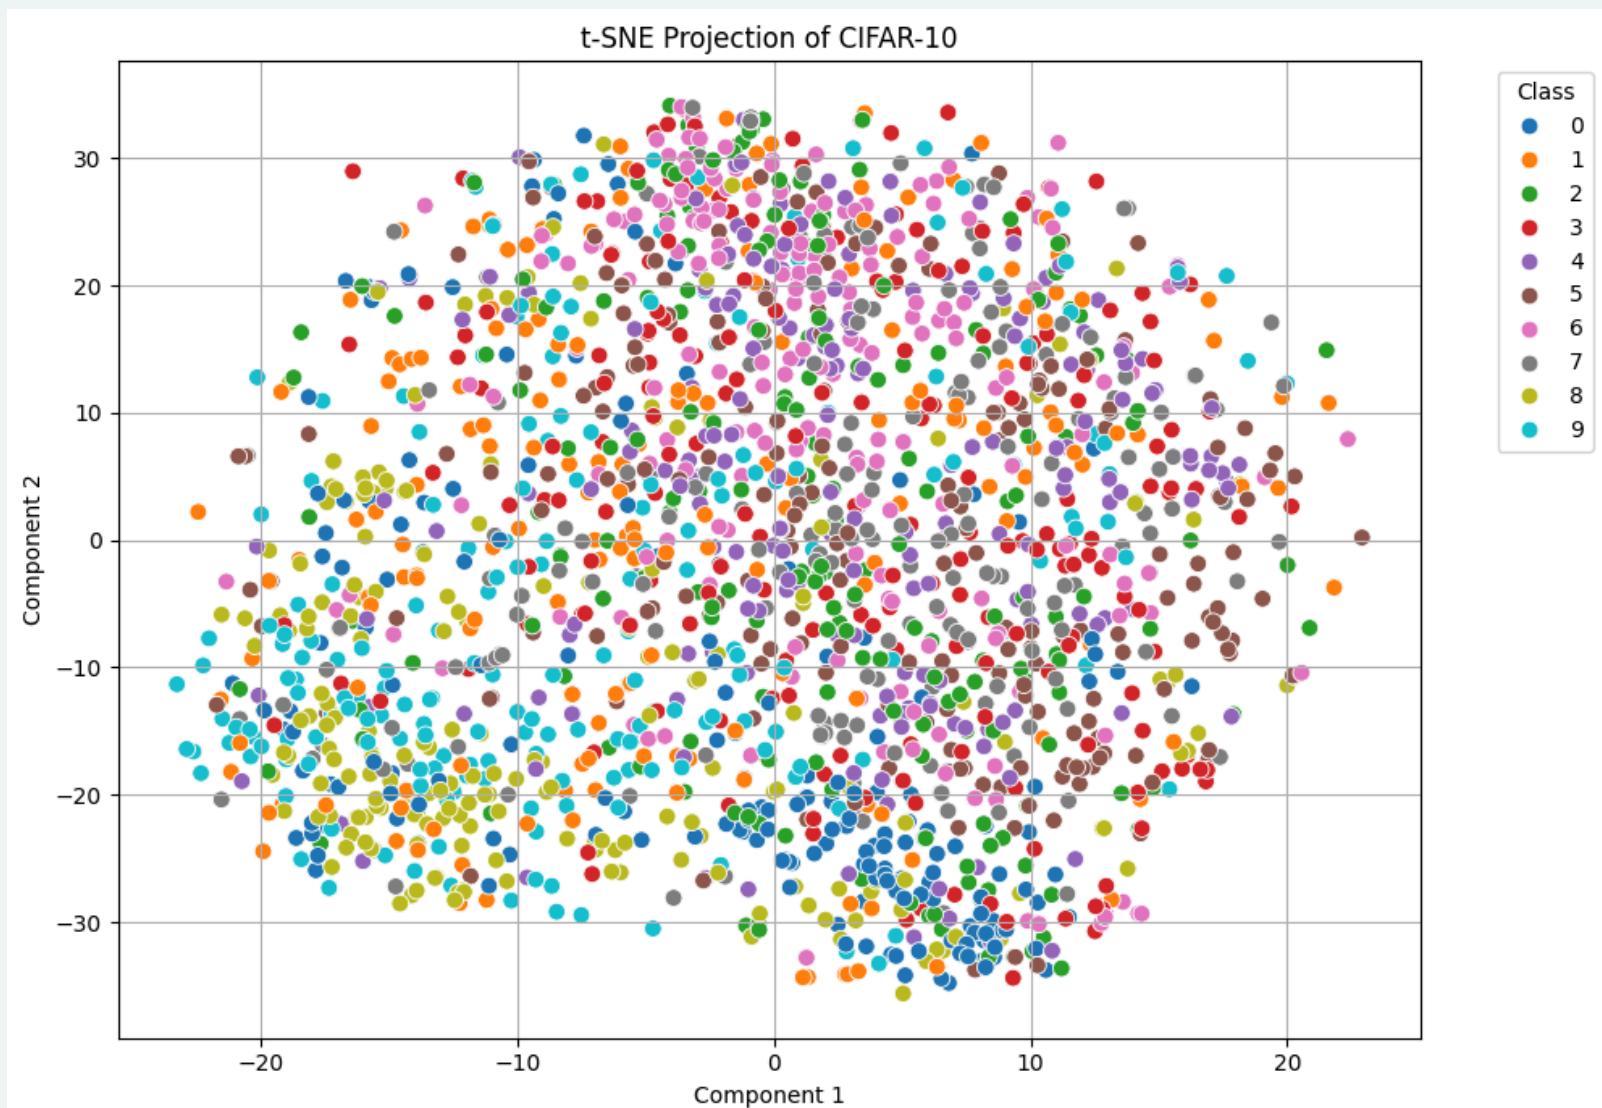
# APPLIED STRATEGIES



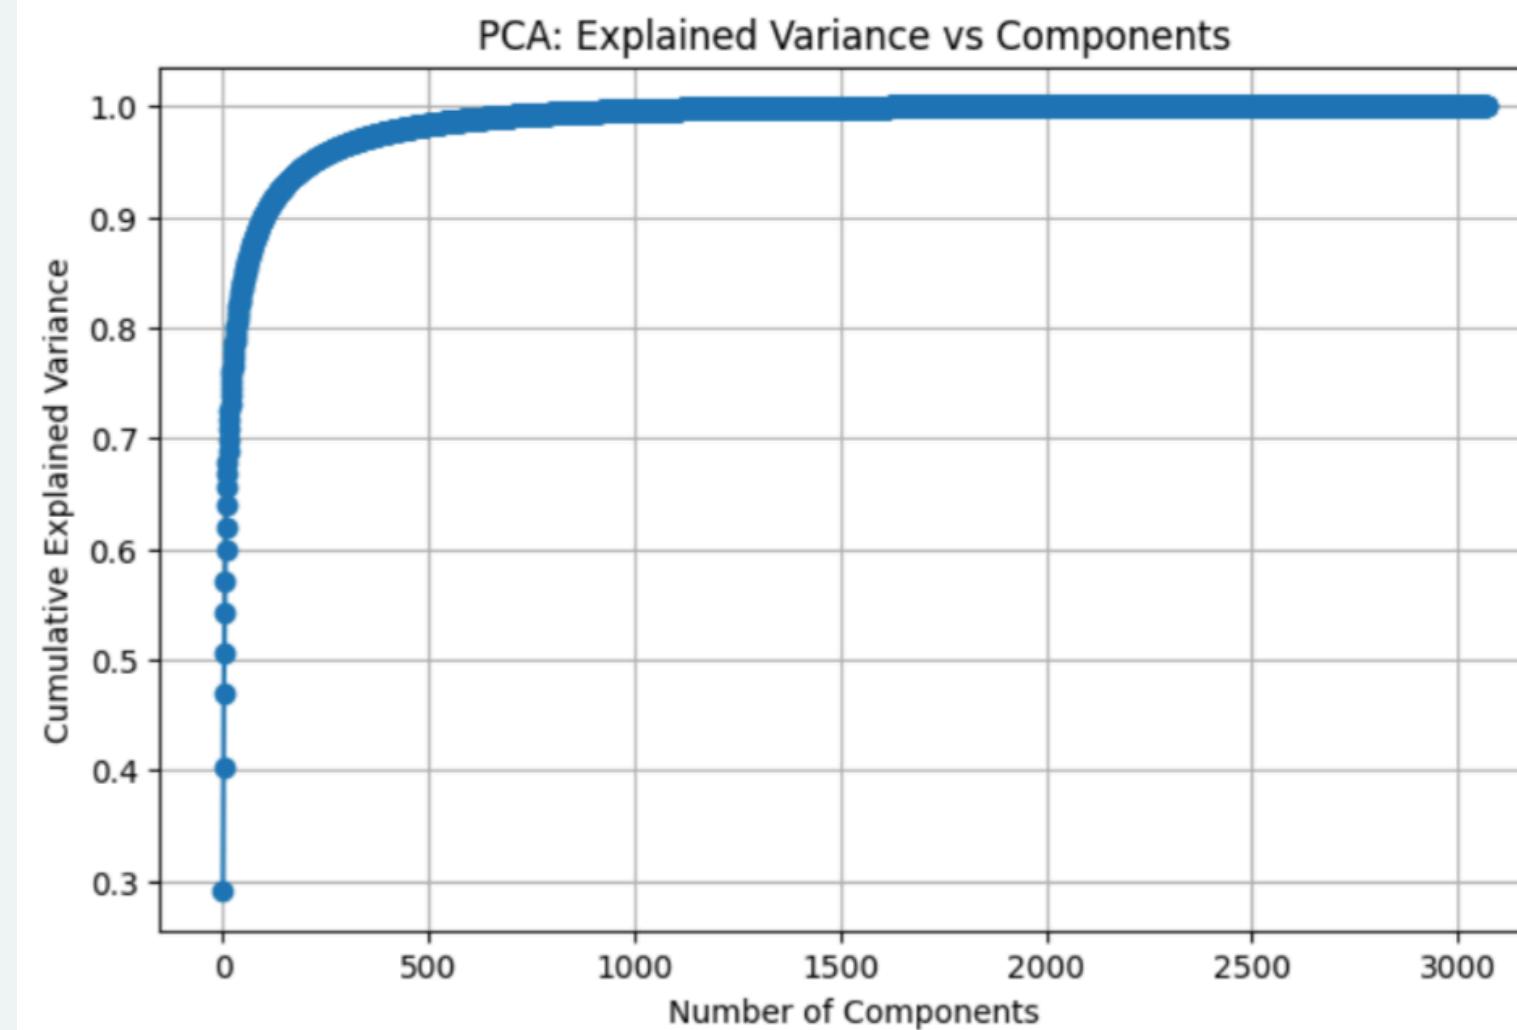
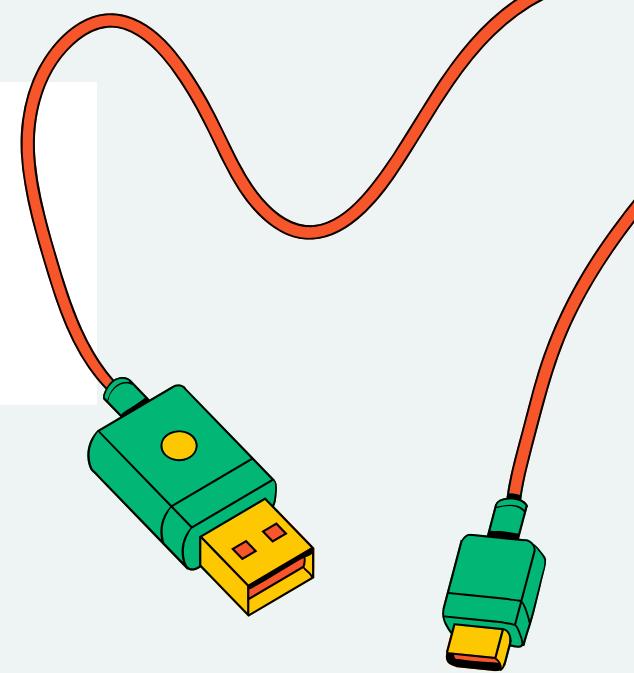
# PREPROCESSING STEPS



O: 'airplane', 1: 'automobile', 2: 'bird', 3: 'cat', 4: 'deer', 5: 'dog', 6: 'frog', 7: 'horse', 8: 'ship', 9: 'truck'



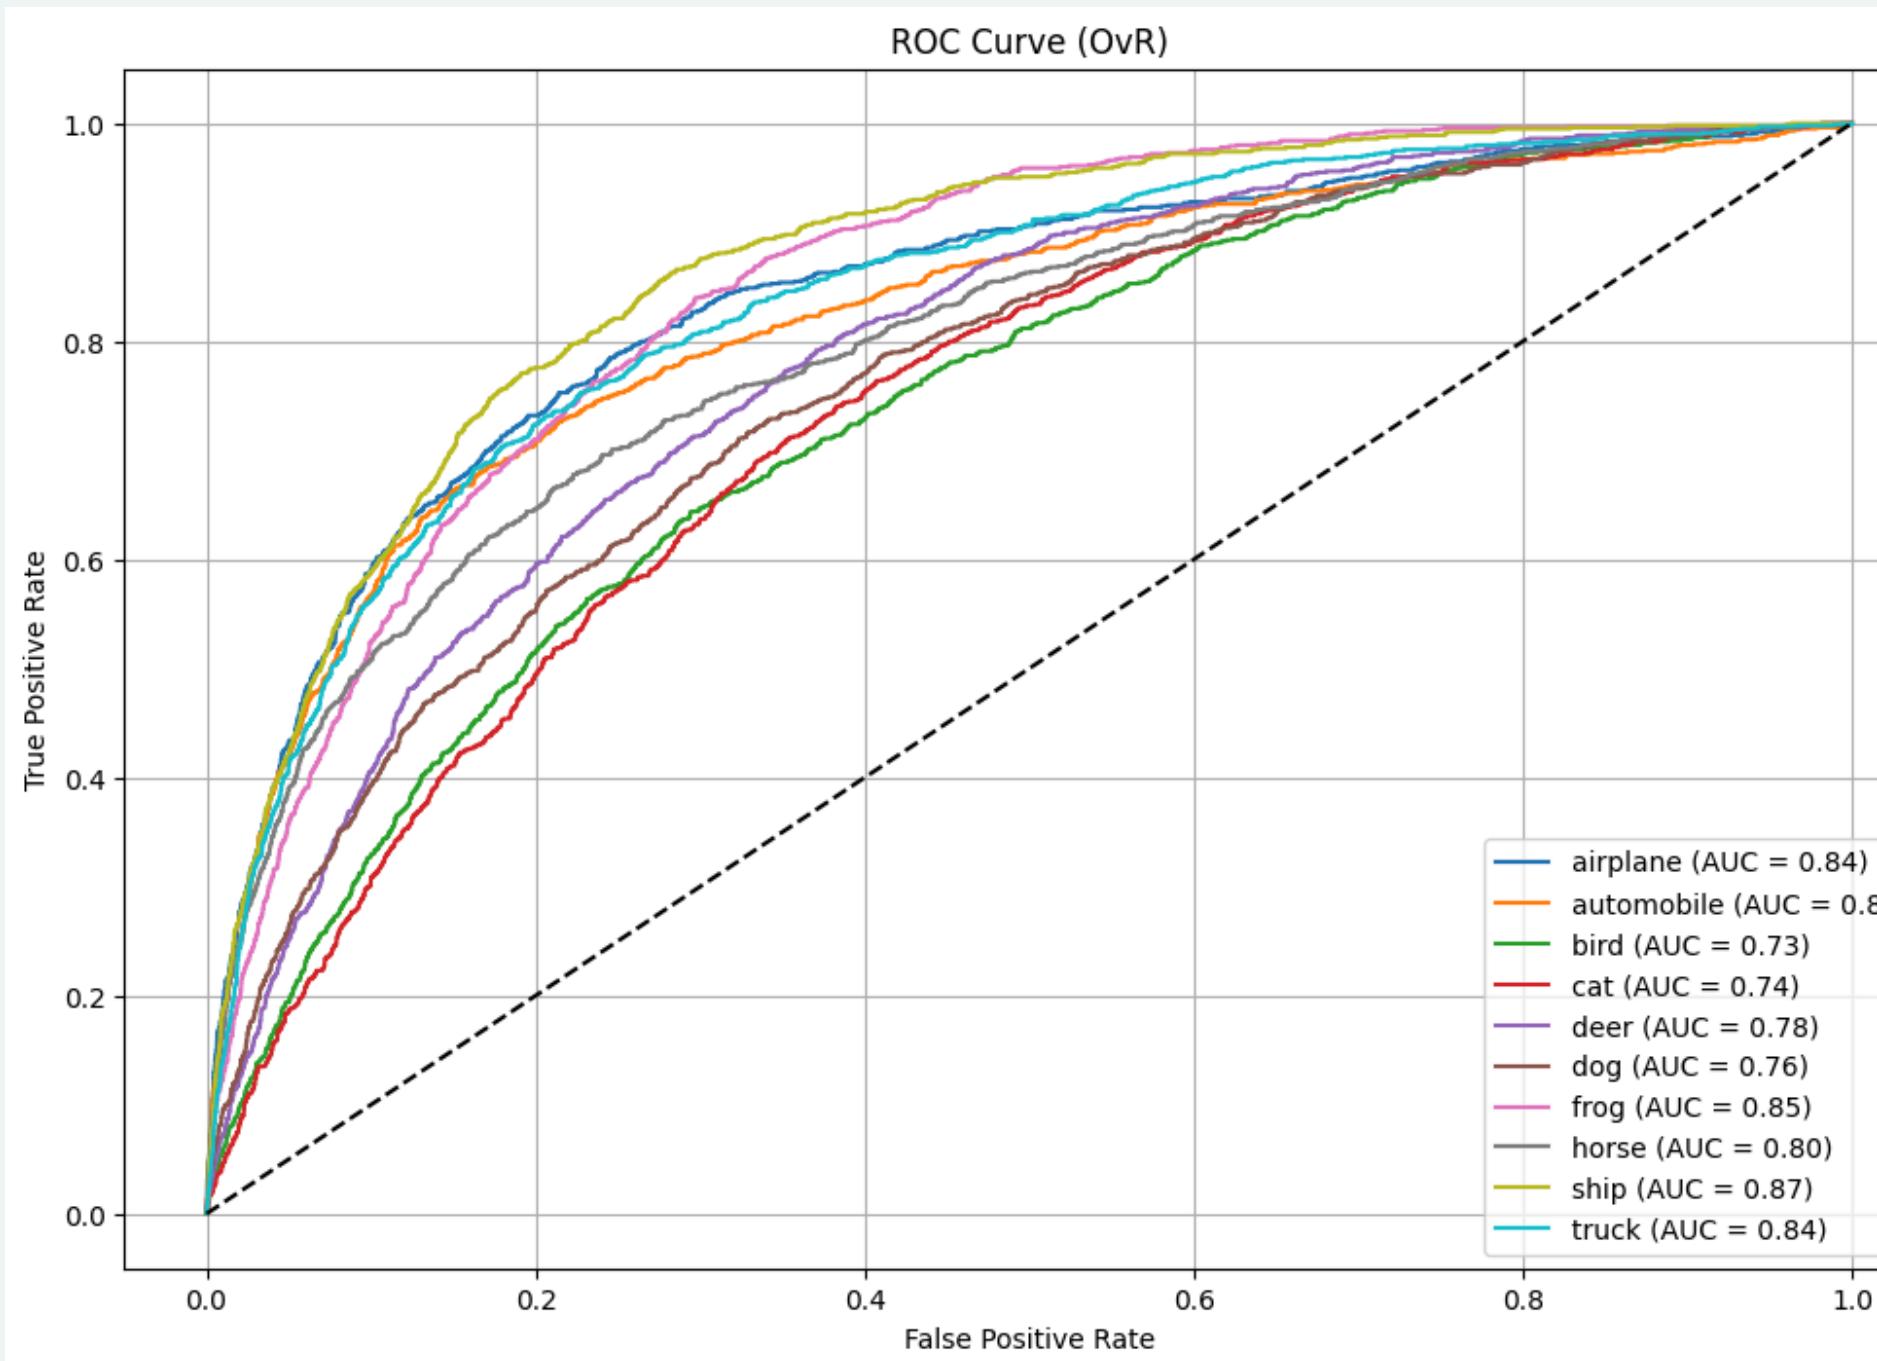
array( [[[178, 176, 189],  
[178, 176, 189],  
[178, 176, 189],



Optimal number of components for 95% variance: 217

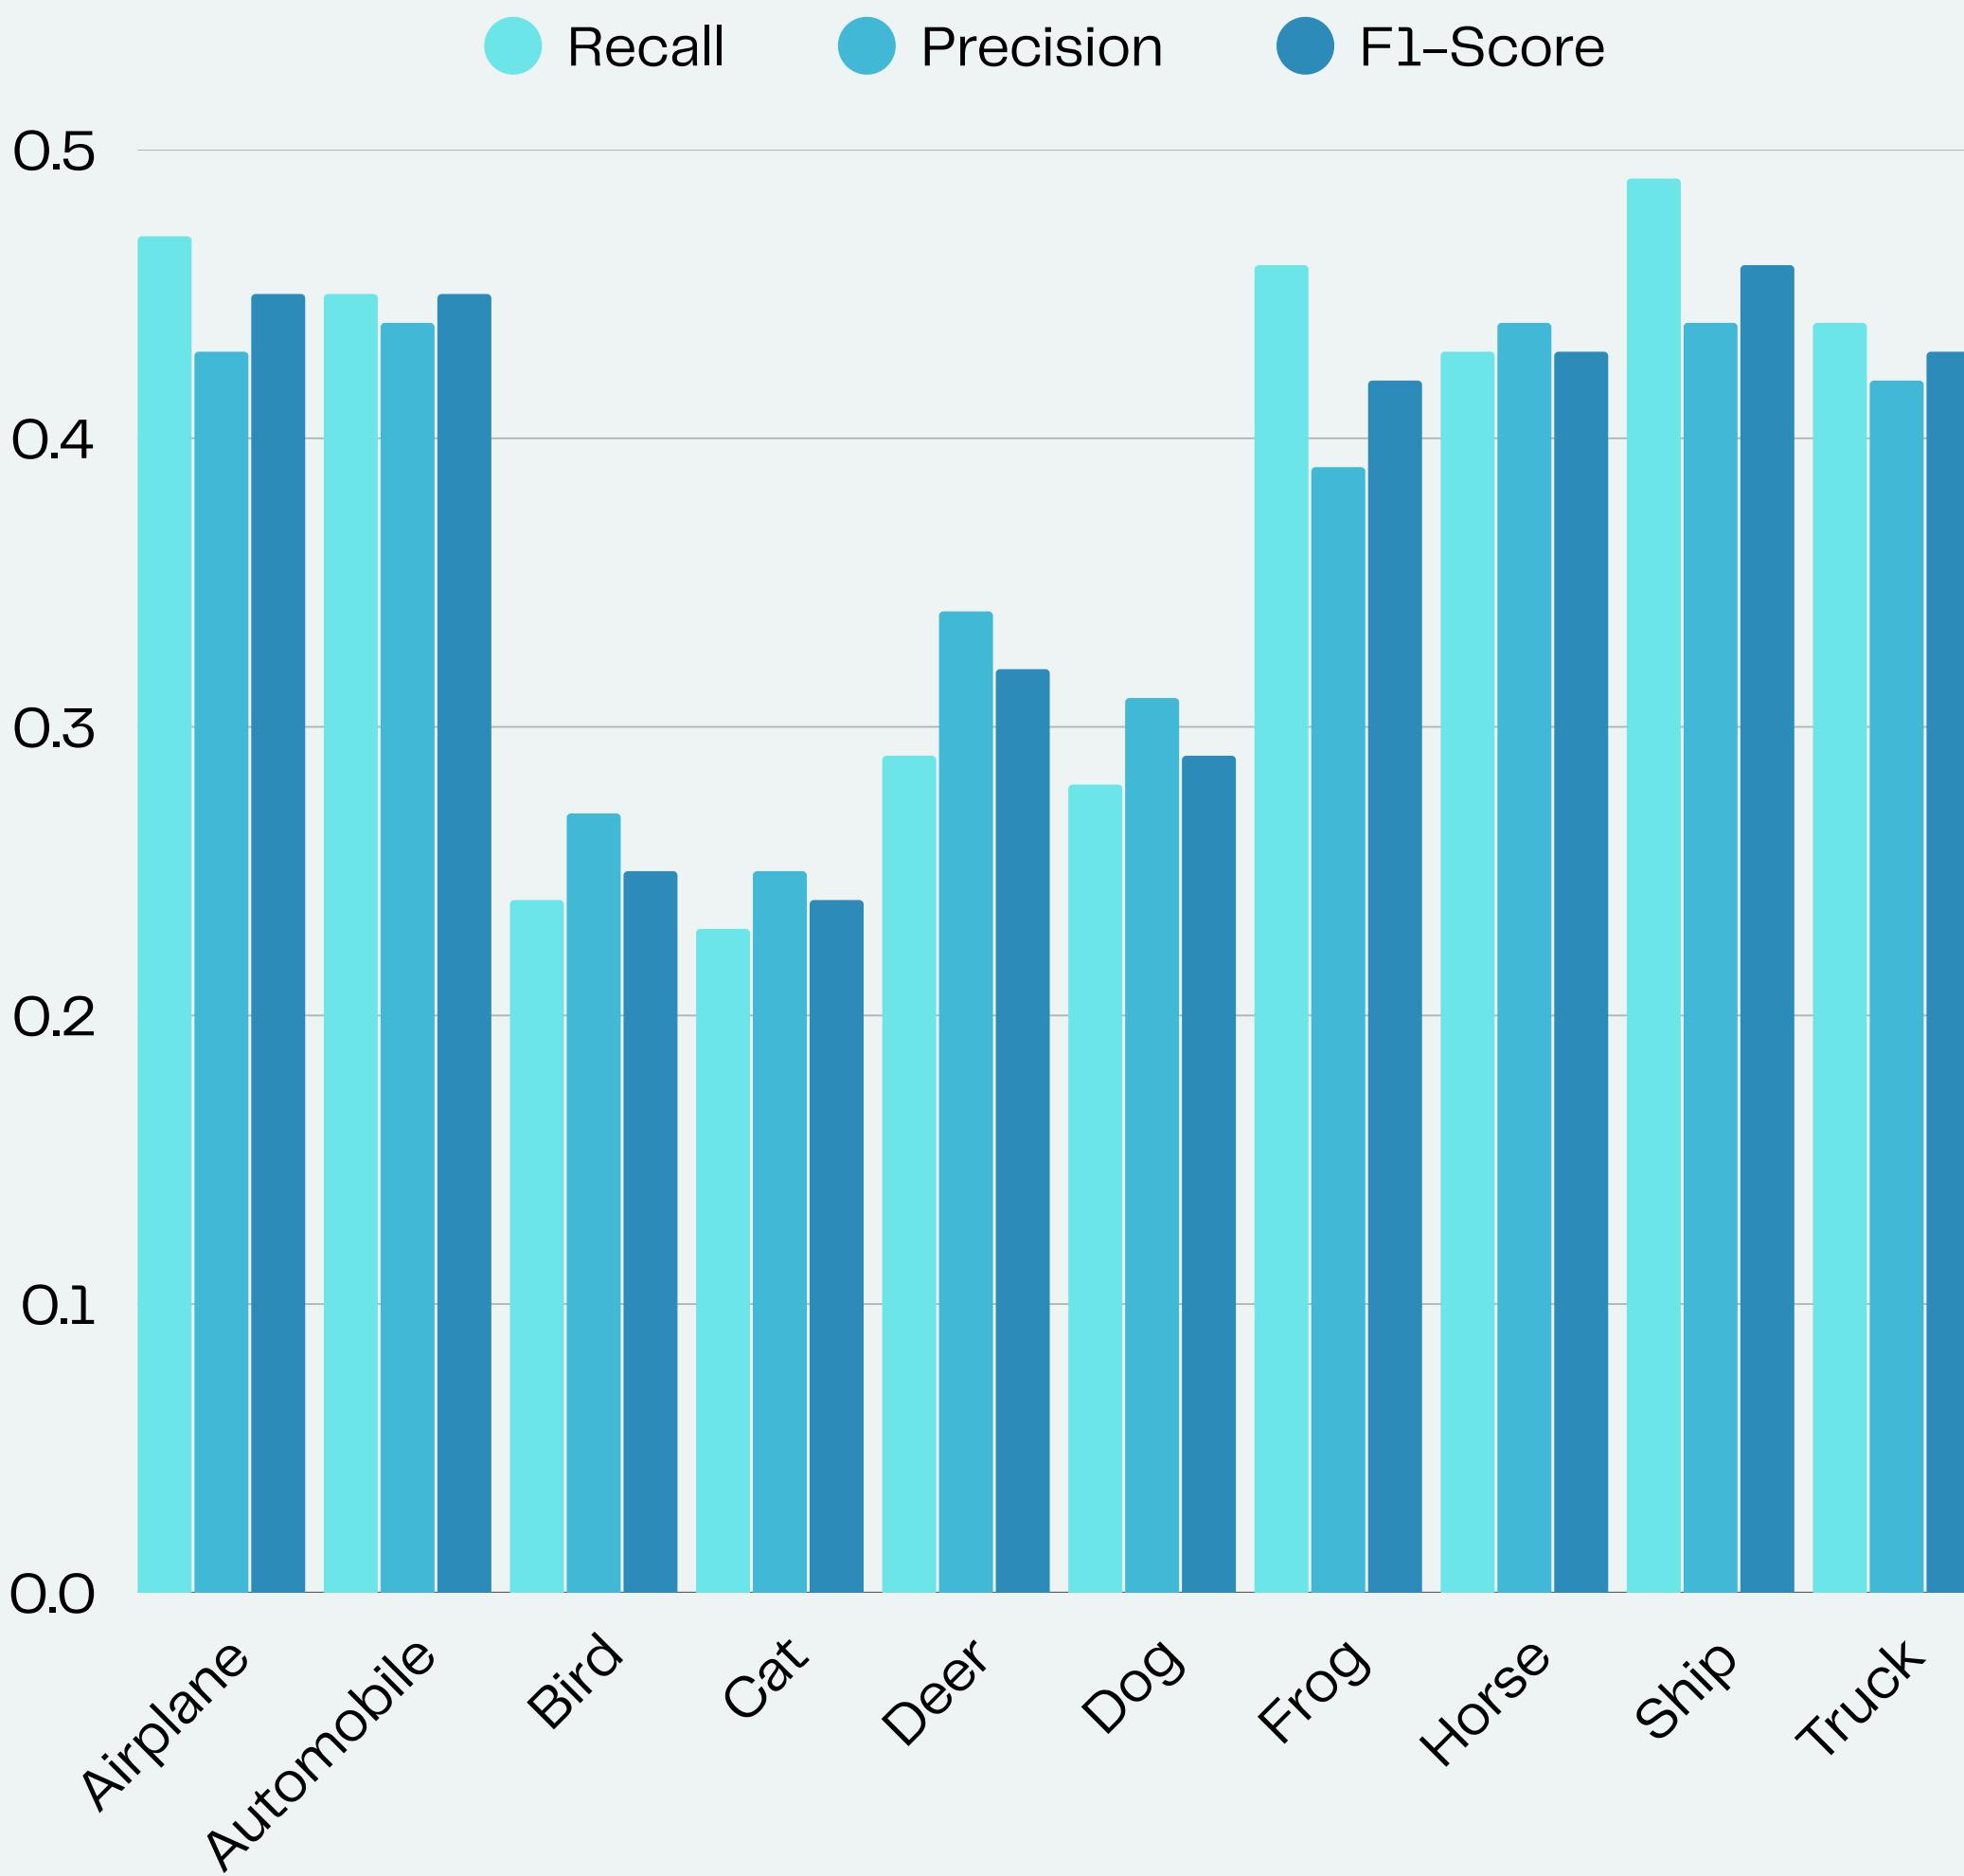
# SOFTMAX

```
clf = LogisticRegression(multi_class='multinomial', solver='lbfgs', max_iter=200)
clf.fit(X_pca, y)
```

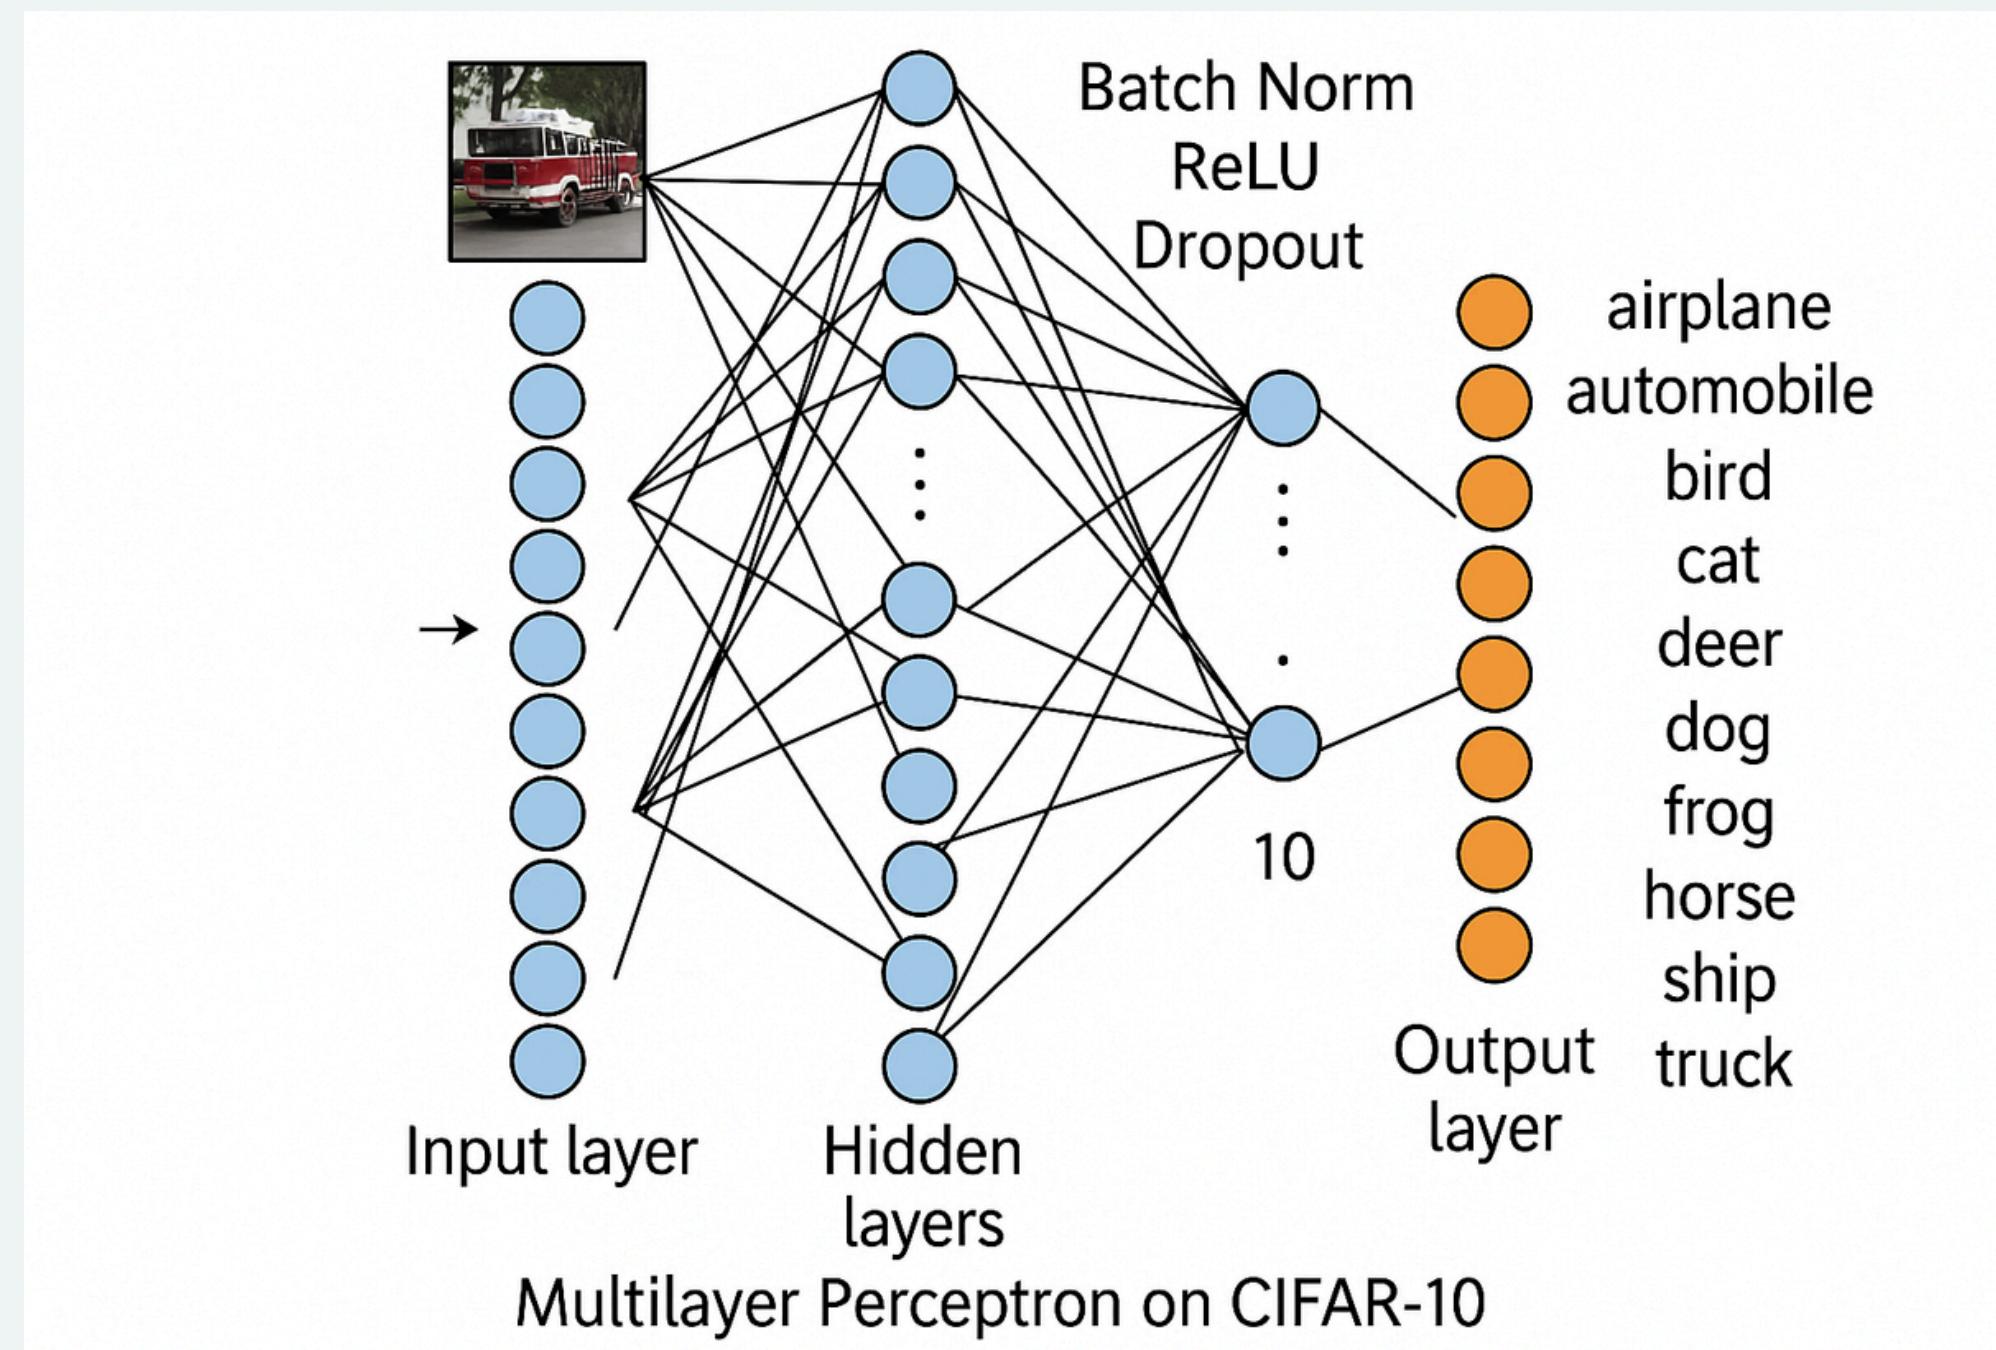
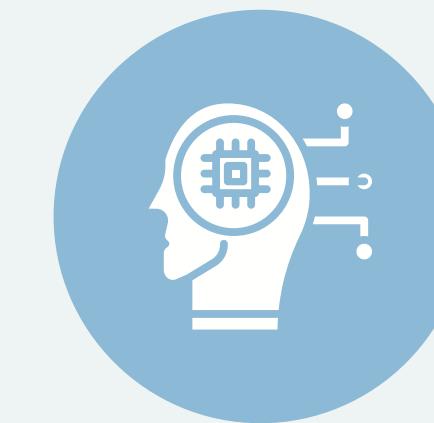


| Metric       | Score |
|--------------|-------|
| Accuracy     | 0.38  |
| Macro Avg    | 0.37  |
| Weighted Avg | 0.37  |

| Class      | Precision | Recall | F1-Score | Support |
|------------|-----------|--------|----------|---------|
| Airplane   | 0.43      | 0.47   | 0.45     | 1000    |
| Automobile | 0.44      | 0.45   | 0.45     | 1000    |
| Bird       | 0.27      | 0.24   | 0.25     | 1000    |
| Cat        | 0.25      | 0.23   | 0.24     | 1000    |
| Deer       | 0.34      | 0.29   | 0.32     | 1000    |
| Dog        | 0.31      | 0.28   | 0.29     | 1000    |
| Frog       | 0.39      | 0.46   | 0.42     | 1000    |
| Horse      | 0.44      | 0.43   | 0.43     | 1000    |
| Ship       | 0.44      | 0.49   | 0.46     | 1000    |
| Truck      | 0.42      | 0.44   | 0.43     | 1000    |



# MULTILAYER PERCEPTRON



- ▶ Training MLP on 500-dim data...
- Epoch 1/80 – acc@1k: 52.30%
- Epoch 2/80 – acc@1k: 59.10%
- Epoch 3/80 – acc@1k: 61.30%
- Epoch 4/80 – acc@1k: 67.20%
- Epoch 5/80 – acc@1k: 67.50%
- Epoch 6/80 – acc@1k: 72.60%
- Epoch 7/80 – acc@1k: 74.10%
- Epoch 8/80 – acc@1k: 73.70%
- Epoch 9/80 – acc@1k: 76.30%

```
# 5) sample predictions
for i in range(5):
    xi = X_test_p[i:i+1]
    pred = mlp.predict(xi)[0]
    print(f"Sample {i} true: {y_test[i]} pred: {pred}")
```

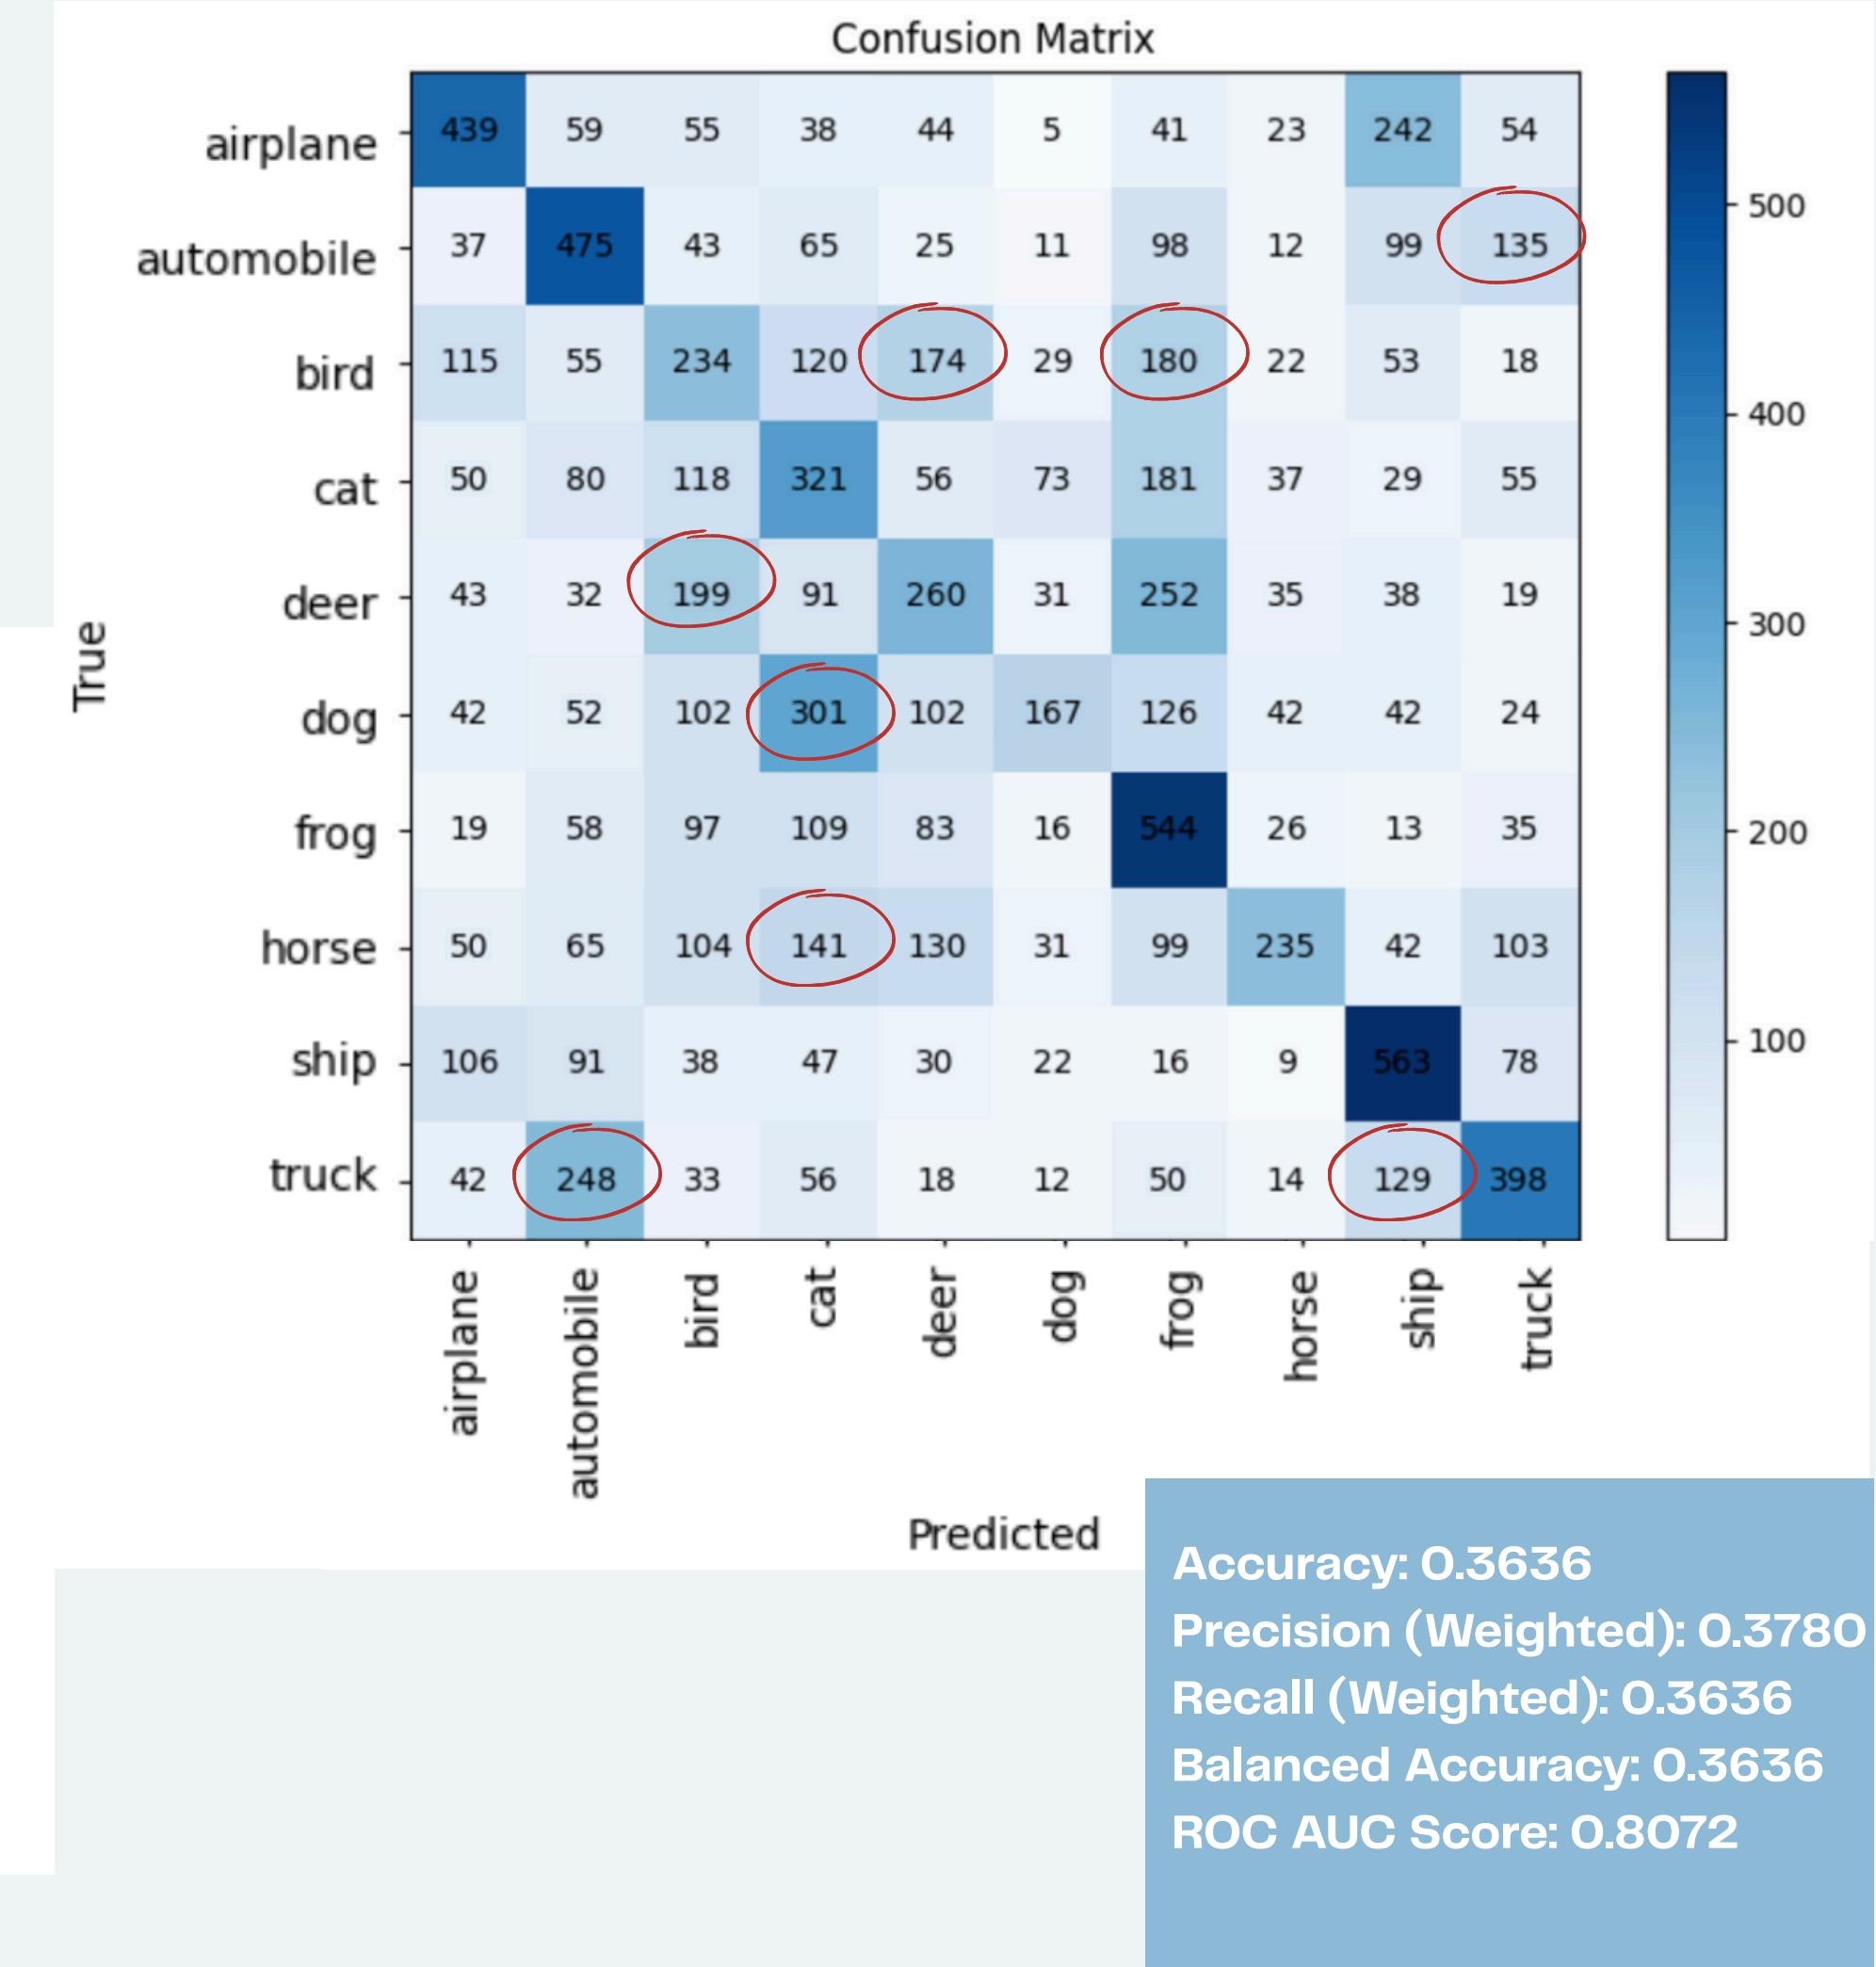
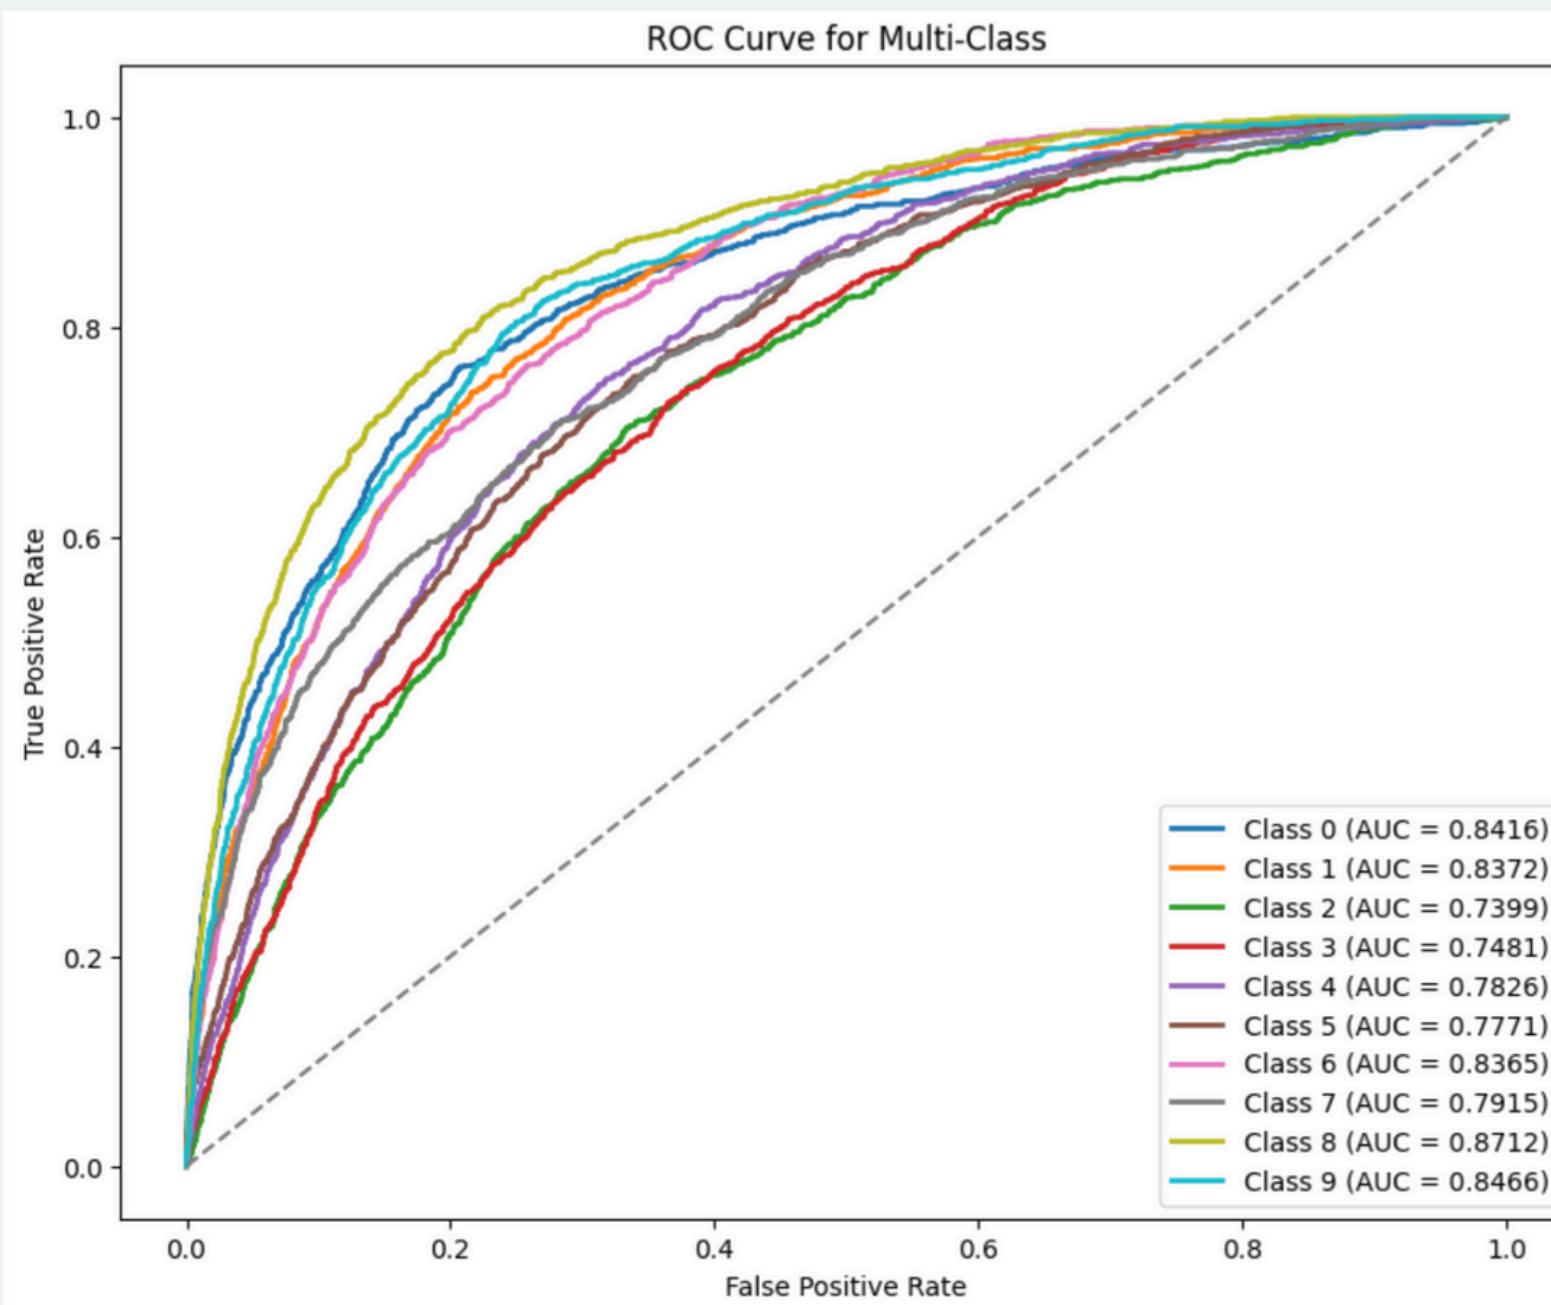
```
Sample 0 true: 3 pred: 3
Sample 1 true: 8 pred: 8
Sample 2 true: 8 pred: 8
Sample 3 true: 0 pred: 8
Sample 4 true: 6 pred: 4
```

► Evaluating on test set...
Test accuracy: 46.98%

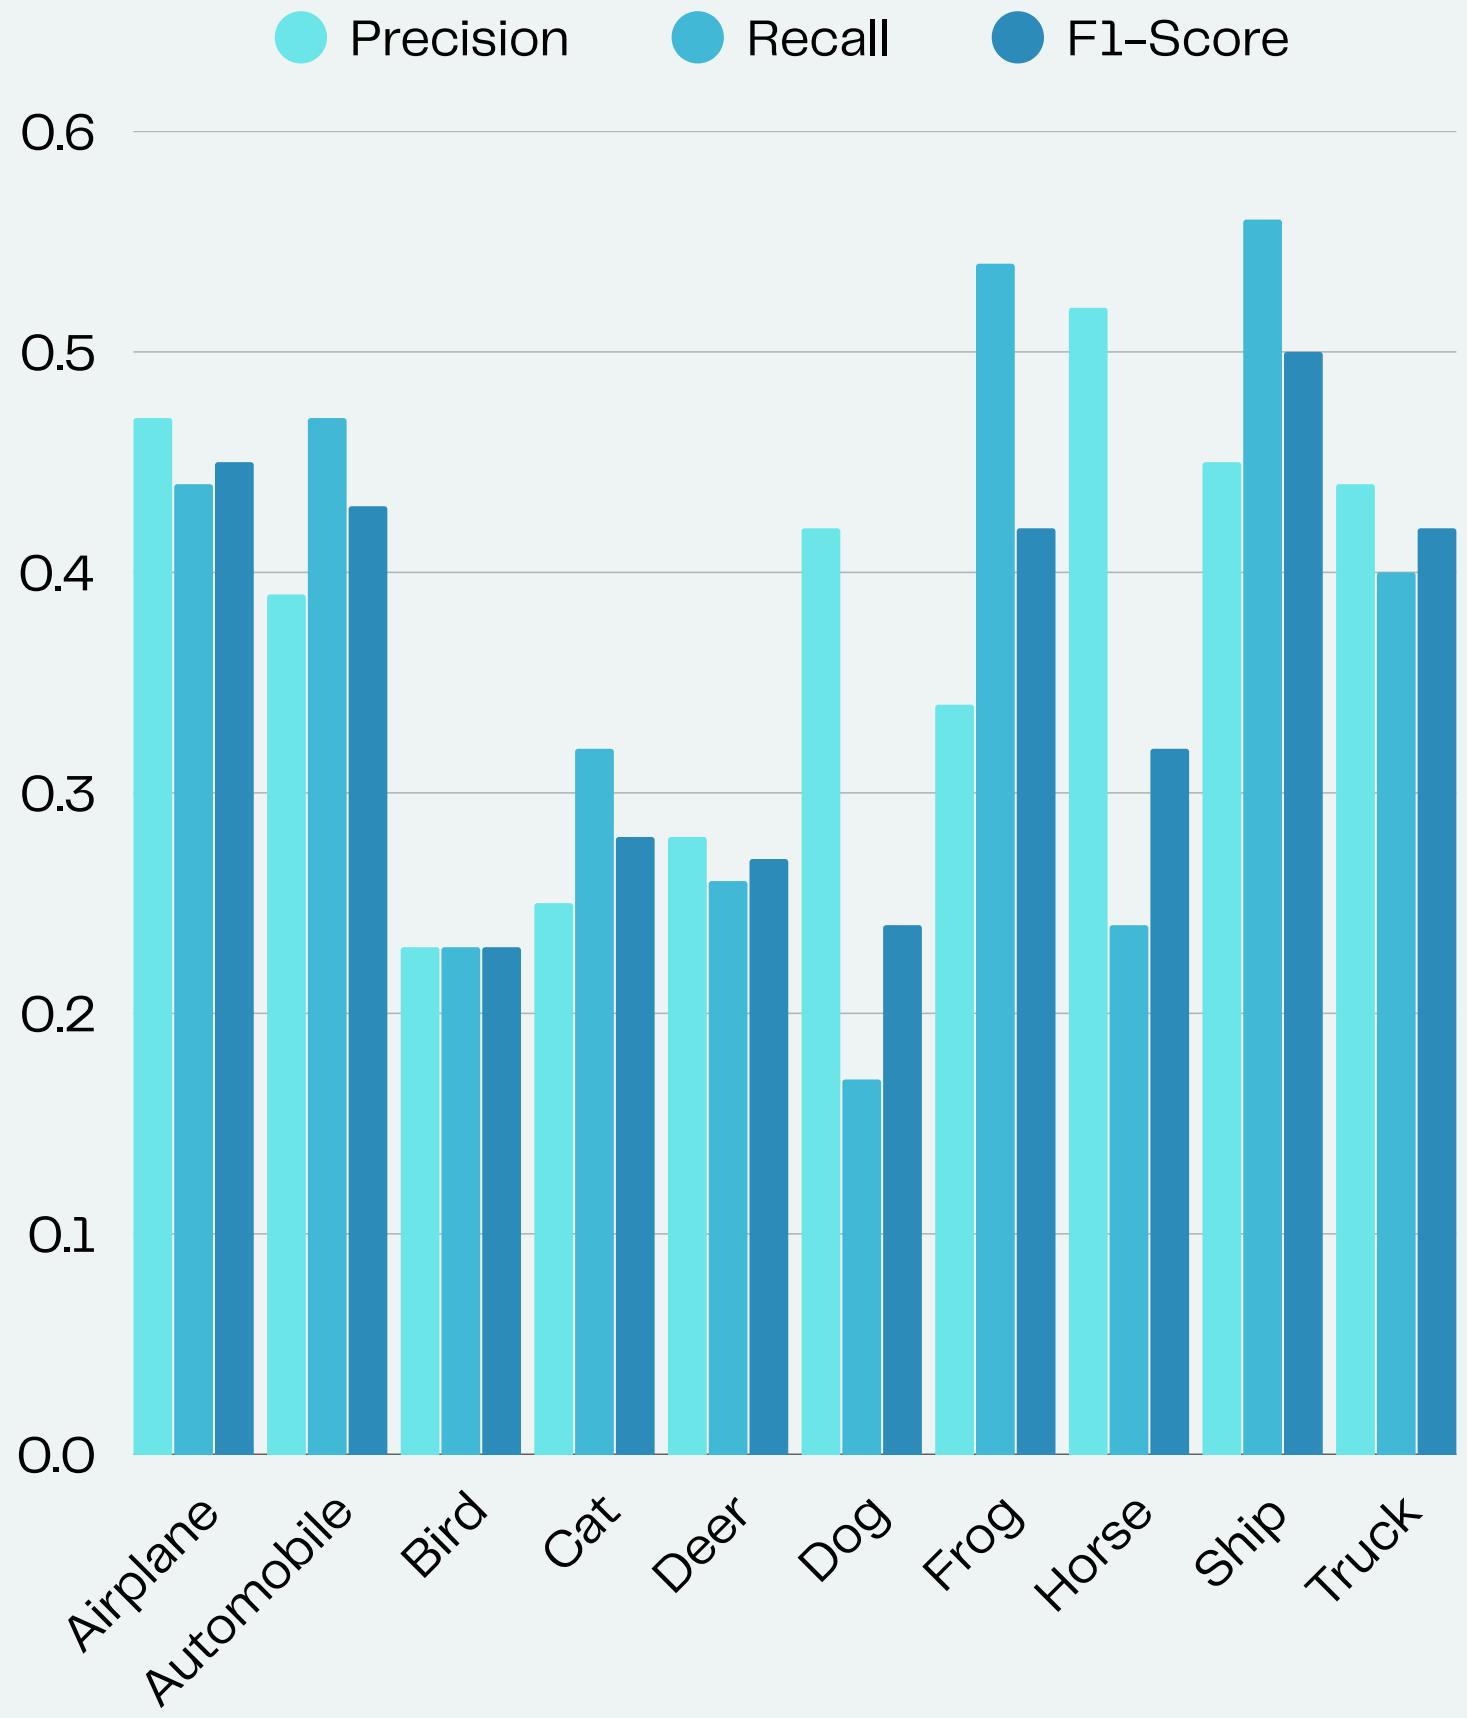
```
class MLP:  
    """  
    Deep fully-connected Perceptron network with:  
        - ReLU activations  
        - Softmax output & cross-entropy loss with label smoothing  
        - Batch normalization on each hidden layer  
        - Dropout on hidden layers  
        - L2 weight decay  
        - Adam optimizer for all parameters  
  
    Parameters  
    -----  
    layer_sizes : list of int  
        Sizes of each layer including input and output, e.g. [300, 512, 256, 128, 10].  
    lr : float, default=1e-3  
        Base learning rate for Adam.  
    dropout_rate : float in [0,1), default=0.5  
        Drop probability on hidden activations.  
    weight_decay : float, default=1e-4  
        L2 regularization coefficient.  
    label_smoothing : float in [0,1), default=0.1  
         $\epsilon$  for smoothing the one-hot targets:  $\hat{y} = (1-\epsilon) \cdot y + \epsilon/K$ .  
    beta1, beta2 : floats, default=(0.9, 0.999)  
        Adam momentum hyperparameters.  
    eps : float, default=1e-8  
        Adam & batch-norm stability constant.  
    """
```

► Evaluating on test set...  
Test accuracy: 58.56%

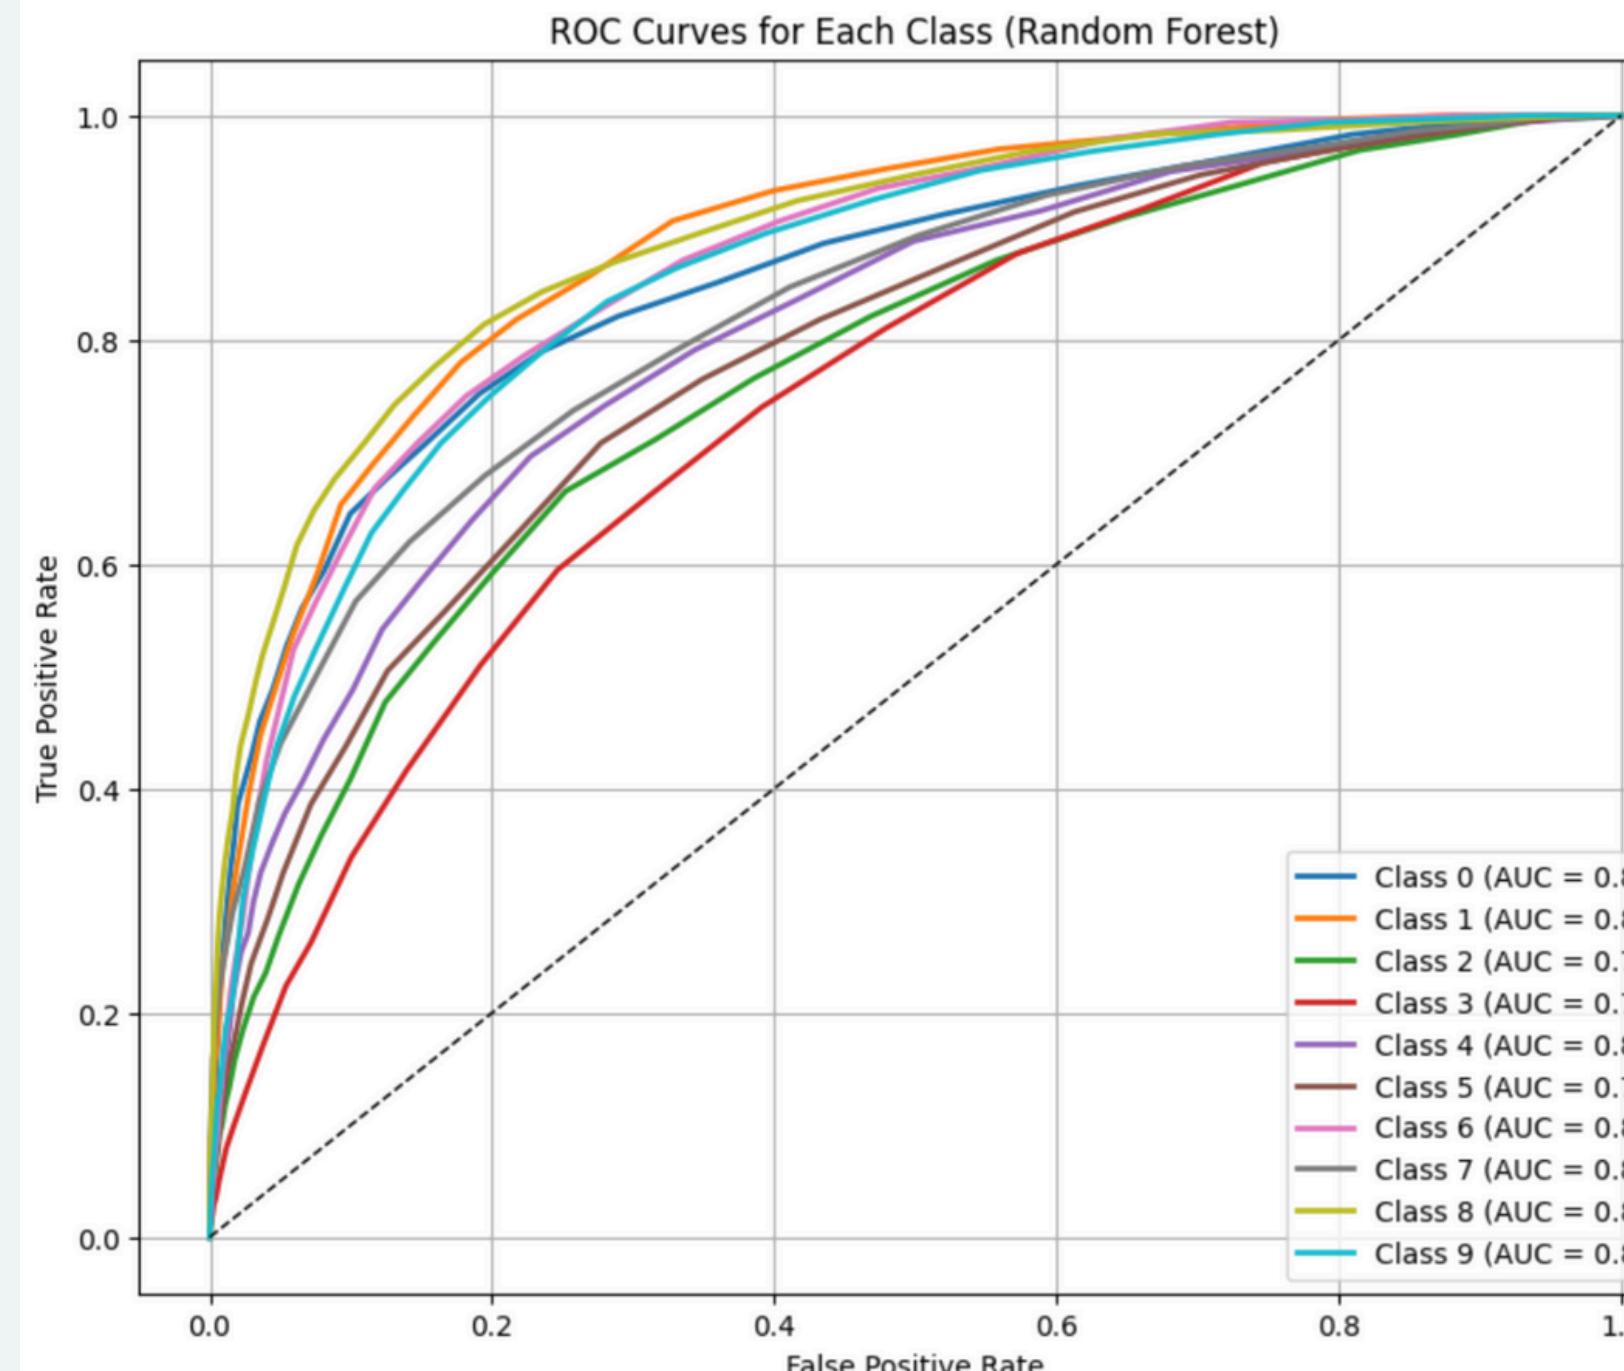
# SVM



| Class      | Precision | Recall | F1-Score | Support |
|------------|-----------|--------|----------|---------|
| Airplane   | 0.47      | 0.44   | 0.45     | 1000    |
| Automobile | 0.39      | 0.47   | 0.43     | 1000    |
| Bird       | 0.23      | 0.23   | 0.23     | 1000    |
| Cat        | 0.25      | 0.32   | 0.28     | 1000    |
| Deer       | 0.28      | 0.26   | 0.27     | 1000    |
| Dog        | 0.42      | 0.17   | 0.24     | 1000    |
| Frog       | 0.34      | 0.54   | 0.42     | 1000    |
| Horse      | 0.52      | 0.24   | 0.32     | 1000    |
| Ship       | 0.45      | 0.56   | 0.5      | 1000    |
| Truck      | 0.44      | 0.4    | 0.42     | 1000    |



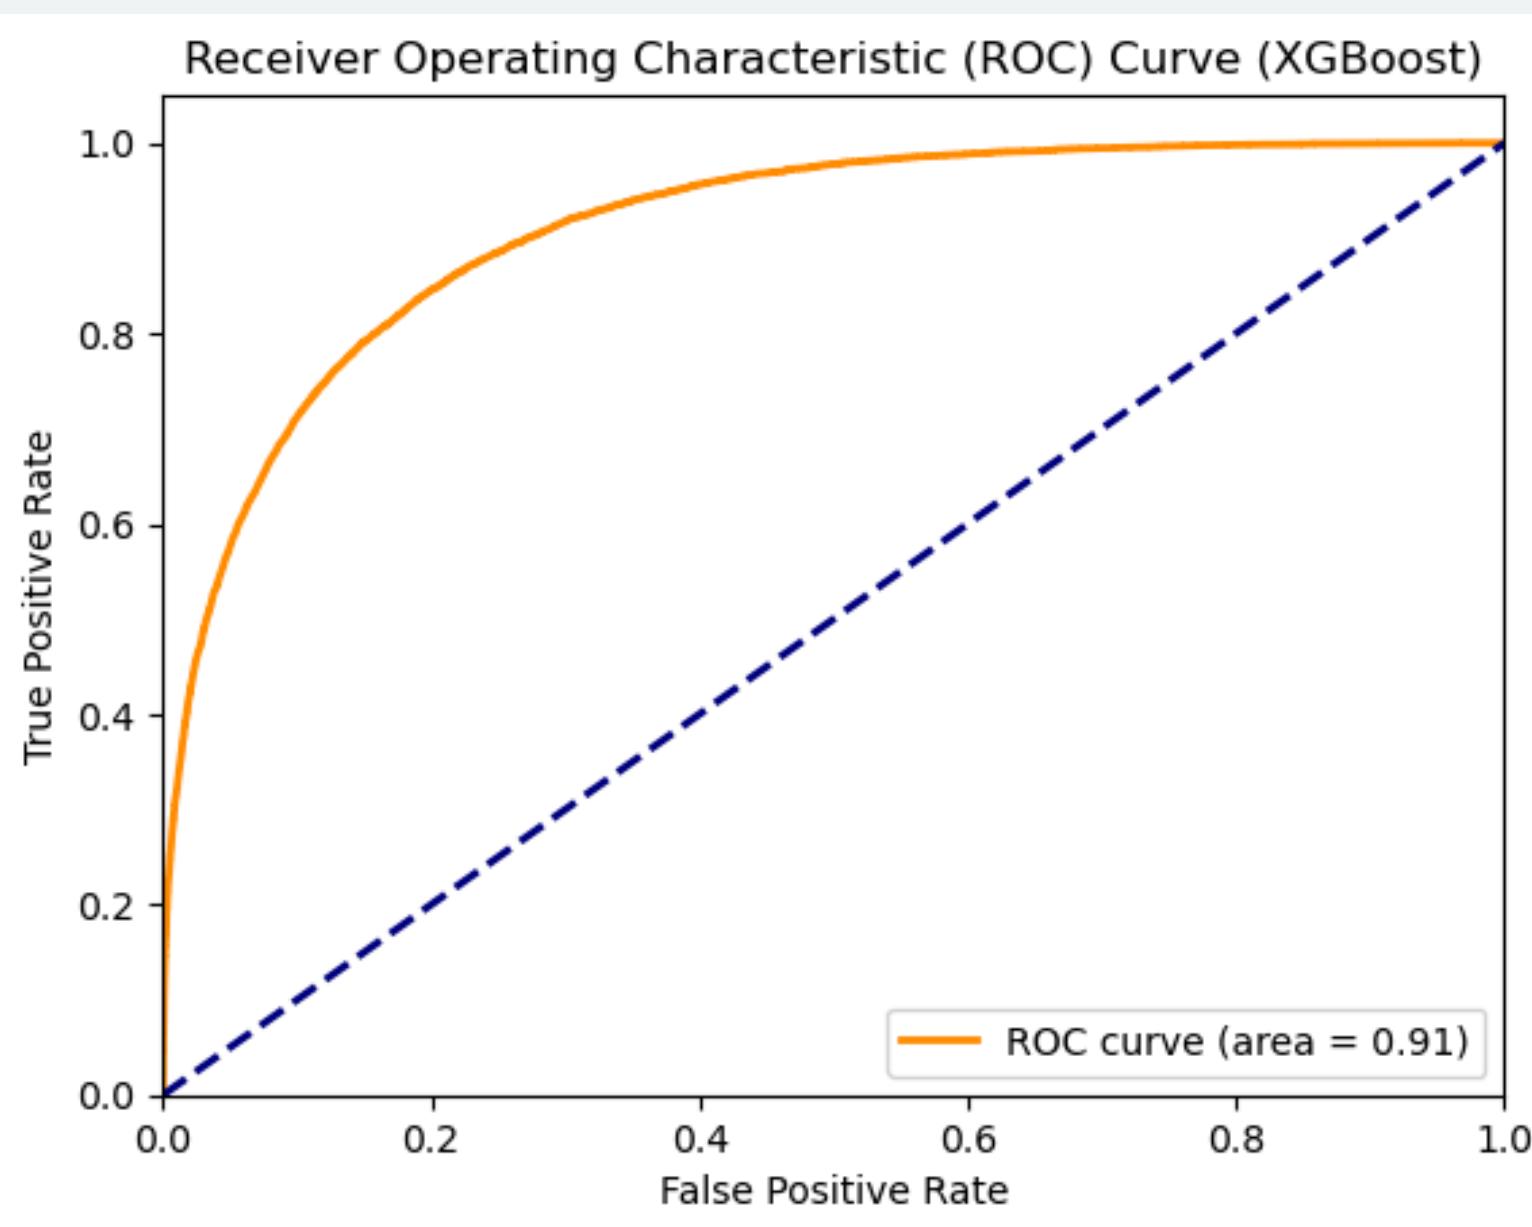
# RANDOM FOREST



| Confusion Matrix |          |            |      |     |      |     |      |       |      |       |
|------------------|----------|------------|------|-----|------|-----|------|-------|------|-------|
|                  | airplane | automobile | bird | cat | deer | dog | frog | horse | ship | truck |
| airplane         | 534      | 63         | 49   | 29  | 30   | 26  | 22   | 31    | 162  | 54    |
| automobile       | 39       | 559        | 19   | 43  | 16   | 27  | 30   | 37    | 61   | 169   |
| bird             | 111      | 35         | 289  | 89  | 156  | 79  | 104  | 64    | 43   | 30    |
| cat              | 51       | 50         | 100  | 255 | 55   | 207 | 113  | 63    | 45   | 61    |
| deer             | 61       | 21         | 151  | 74  | 394  | 42  | 136  | 59    | 37   | 25    |
| dog              | 48       | 44         | 89   | 179 | 66   | 362 | 76   | 68    | 36   | 32    |
| frog             | 15       | 45         | 87   | 74  | 101  | 67  | 526  | 35    | 18   | 32    |
| horse            | 50       | 55         | 46   | 87  | 75   | 94  | 57   | 413   | 32   | 91    |
| ship             | 106      | 92         | 17   | 16  | 14   | 36  | 12   | 18    | 616  | 73    |
| truck            | 58       | 211        | 18   | 35  | 20   | 31  | 25   | 41    | 87   | 474   |

Accuracy: 0.4422  
 Precision: 0.4368  
 Recall: 0.4422  
 Balanced Accuracy: 0.4422  
 ROC AUC Score: 0.8274

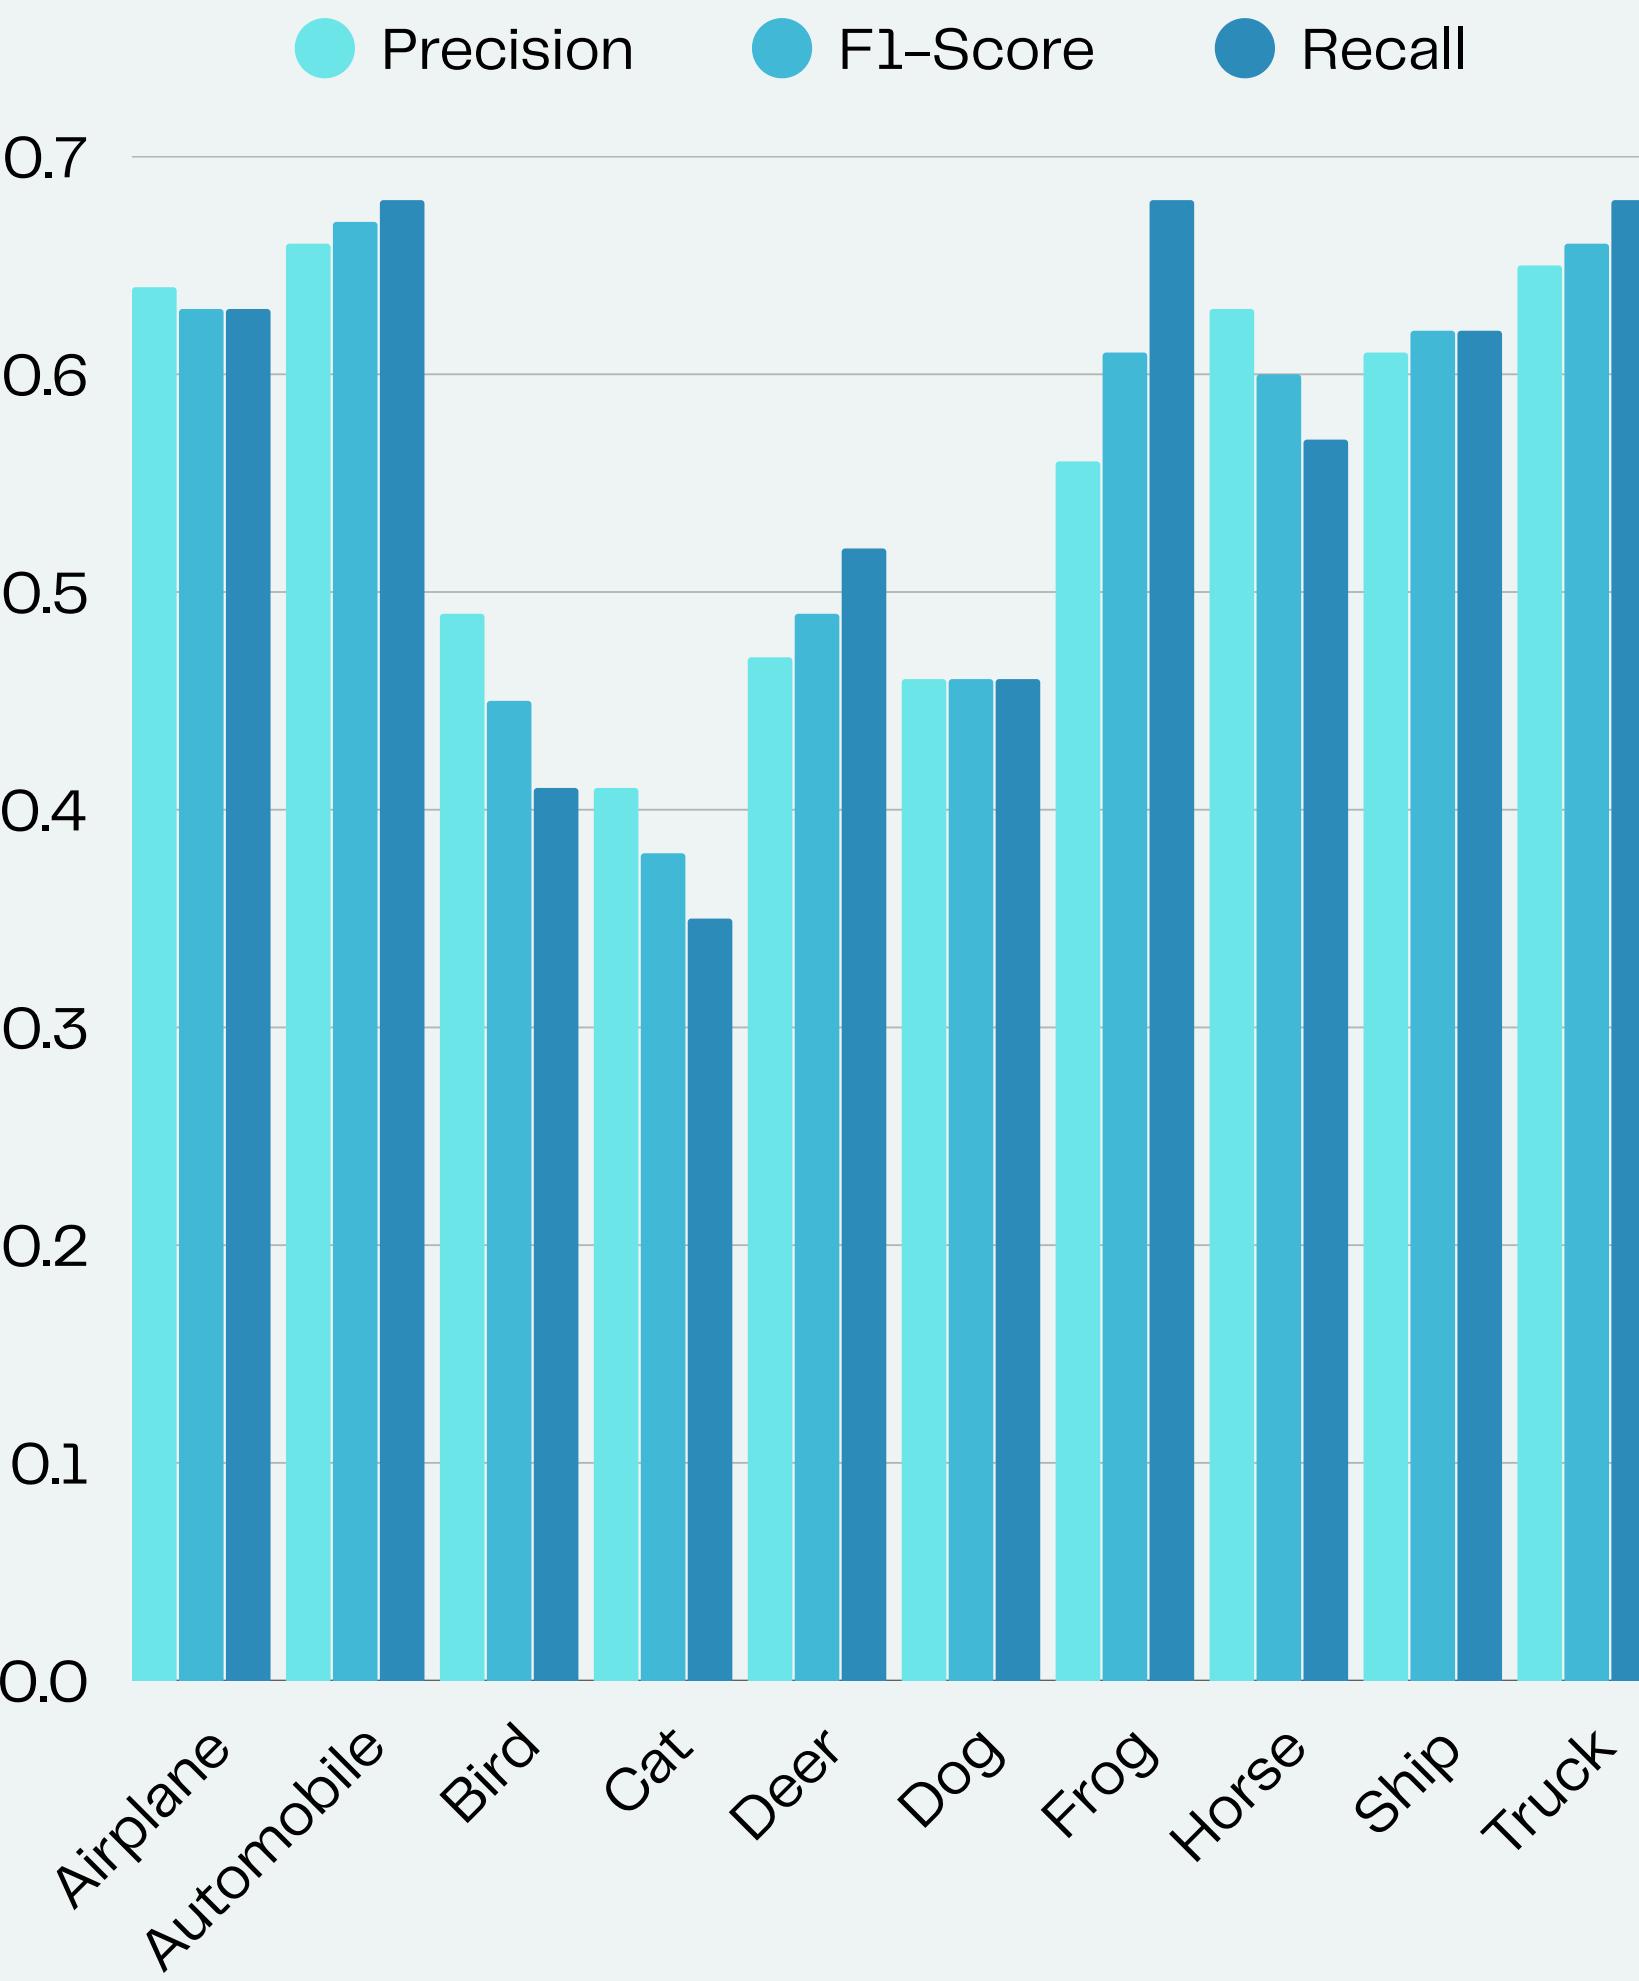
# XGBOOST MODEL WITH HOG FEATURE EXTRACTION



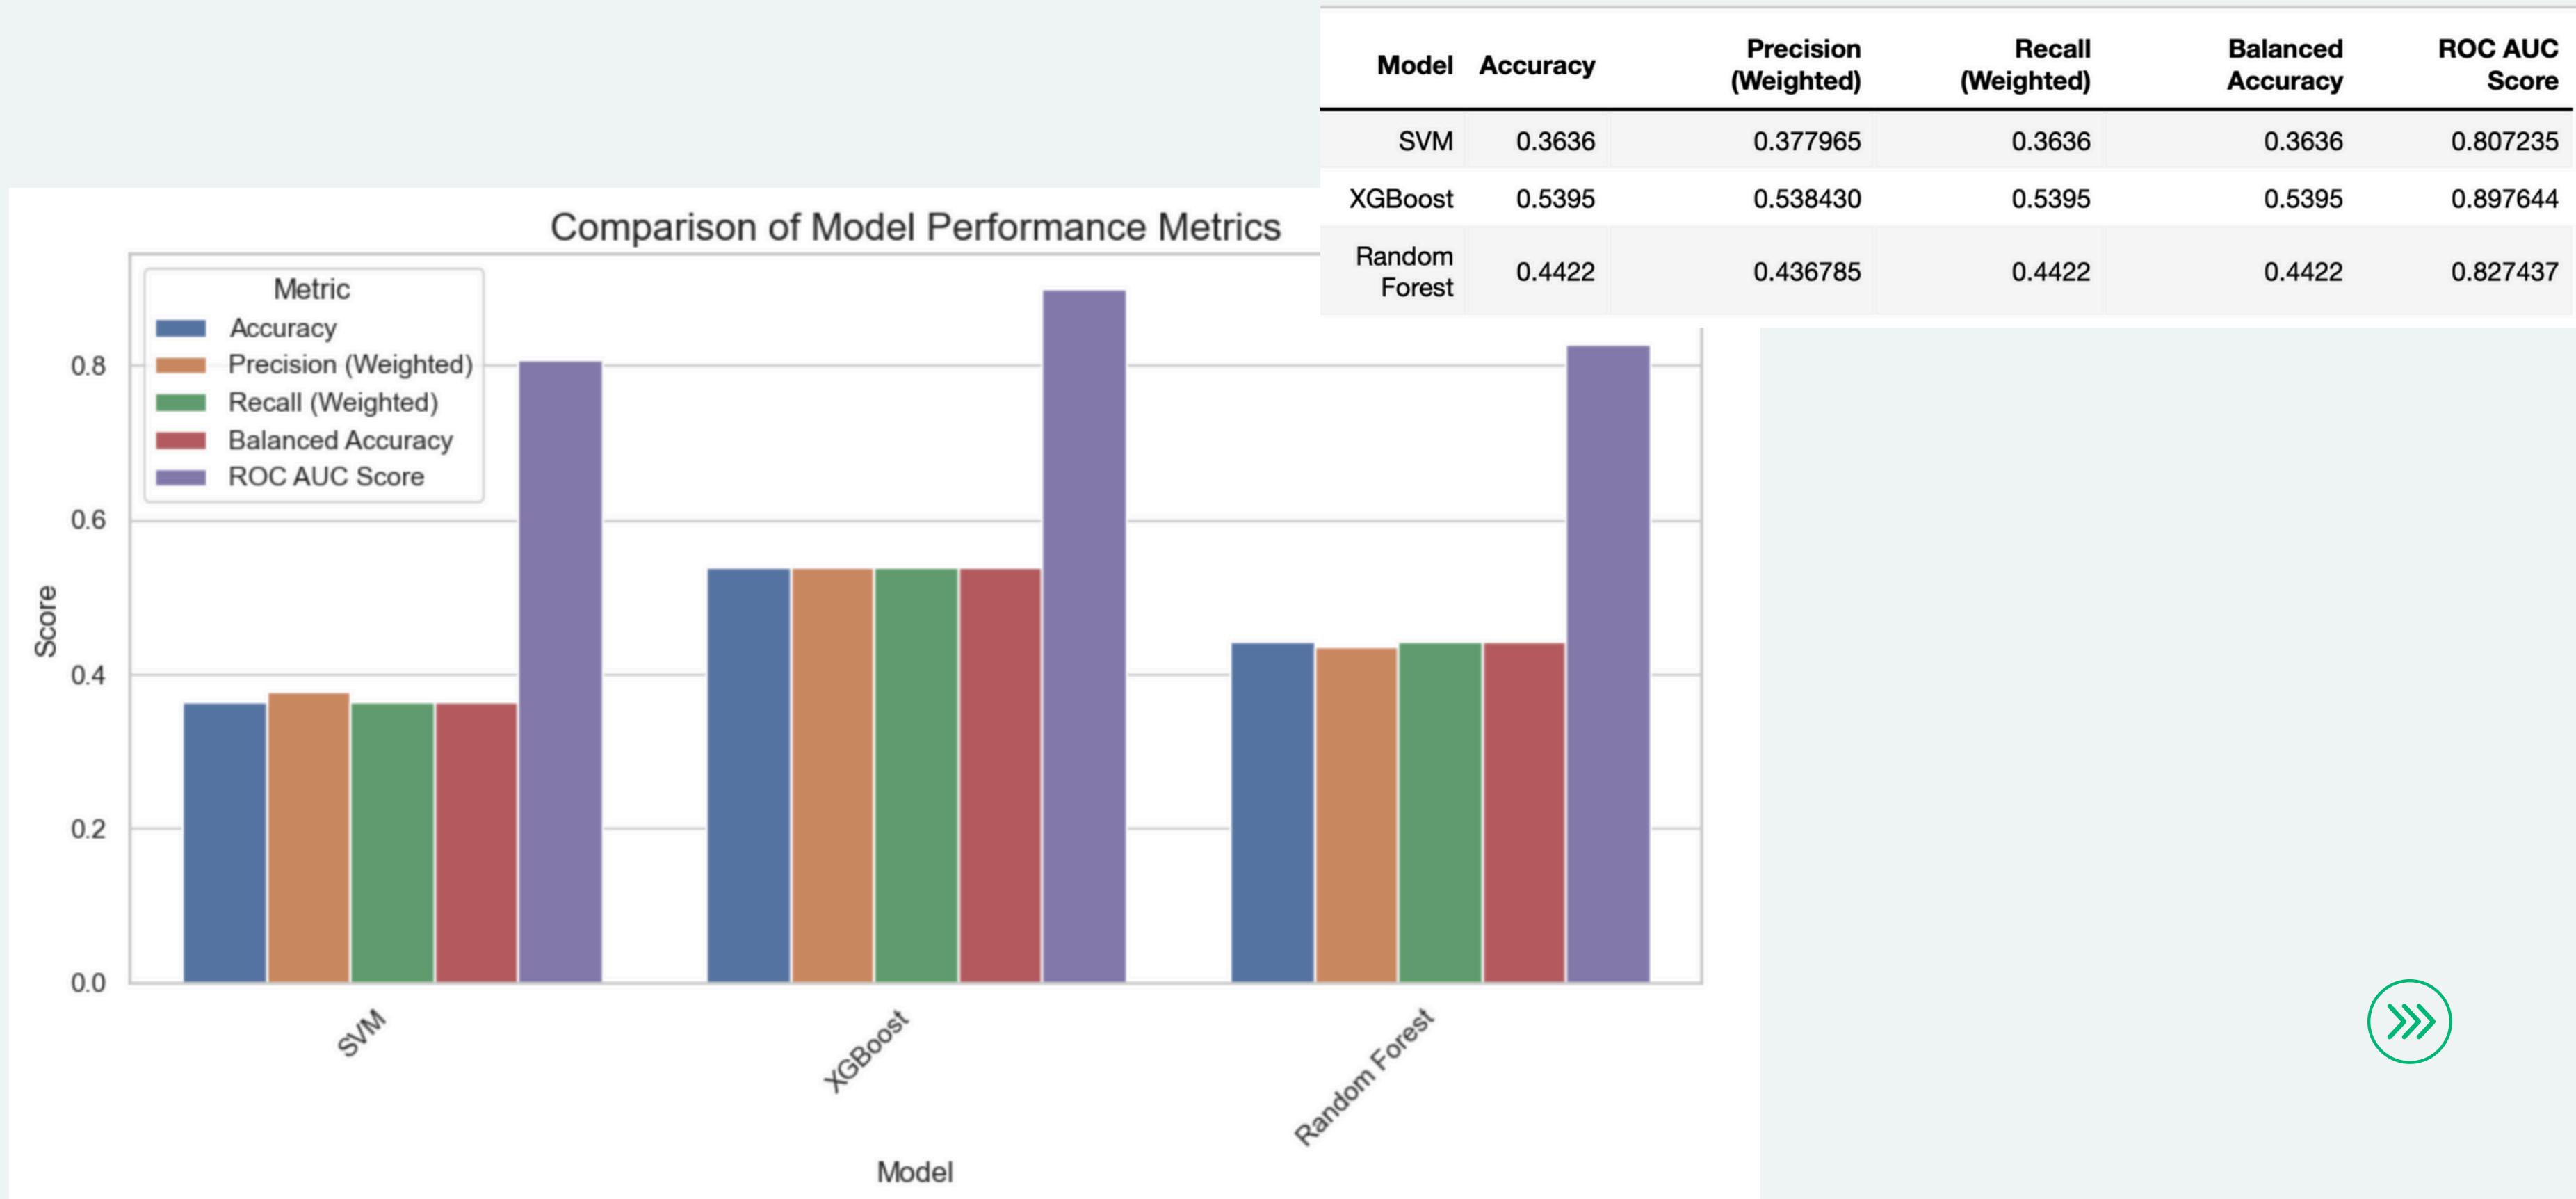
| XGBoost + HOG Confusion Matrix |          |            |      |     |      |     |      |       |      |       |
|--------------------------------|----------|------------|------|-----|------|-----|------|-------|------|-------|
| True                           | airplane | automobile | bird | cat | deer | dog | frog | horse | ship | truck |
| airplane                       | - 628    | 27         | 83   | 18  | 33   | 10  | 23   | 11    | 133  | 34    |
| automobile                     | - 23     | 680        | 8    | 23  | 23   | 9   | 47   | 11    | 83   | 93    |
| bird                           | - 79     | 19         | 410  | 72  | 125  | 112 | 91   | 40    | 35   | 17    |
| cat                            | - 36     | 35         | 74   | 350 | 93   | 176 | 119  | 51    | 20   | 46    |
| deer                           | - 26     | 16         | 56   | 81  | 520  | 61  | 106  | 81    | 30   | 23    |
| dog                            | - 12     | 15         | 58   | 160 | 93   | 456 | 81   | 83    | 15   | 27    |
| frog                           | - 20     | 32         | 43   | 39  | 74   | 67  | 675  | 17    | 17   | 16    |
| horse                          | - 24     | 14         | 63   | 67  | 99   | 75  | 35   | 573   | 13   | 37    |
| ship                           | - 106    | 96         | 27   | 14  | 25   | 11  | 11   | 13    | 625  | 72    |
| truck                          | - 28     | 102        | 16   | 23  | 23   | 22  | 19   | 30    | 61   | 676   |

Accuracy: 0.5593  
Precision : 0.5568  
Recall : 0.5593  
Balanced Accuracy : 0.5593  
ROC AUC Score : 0.9043

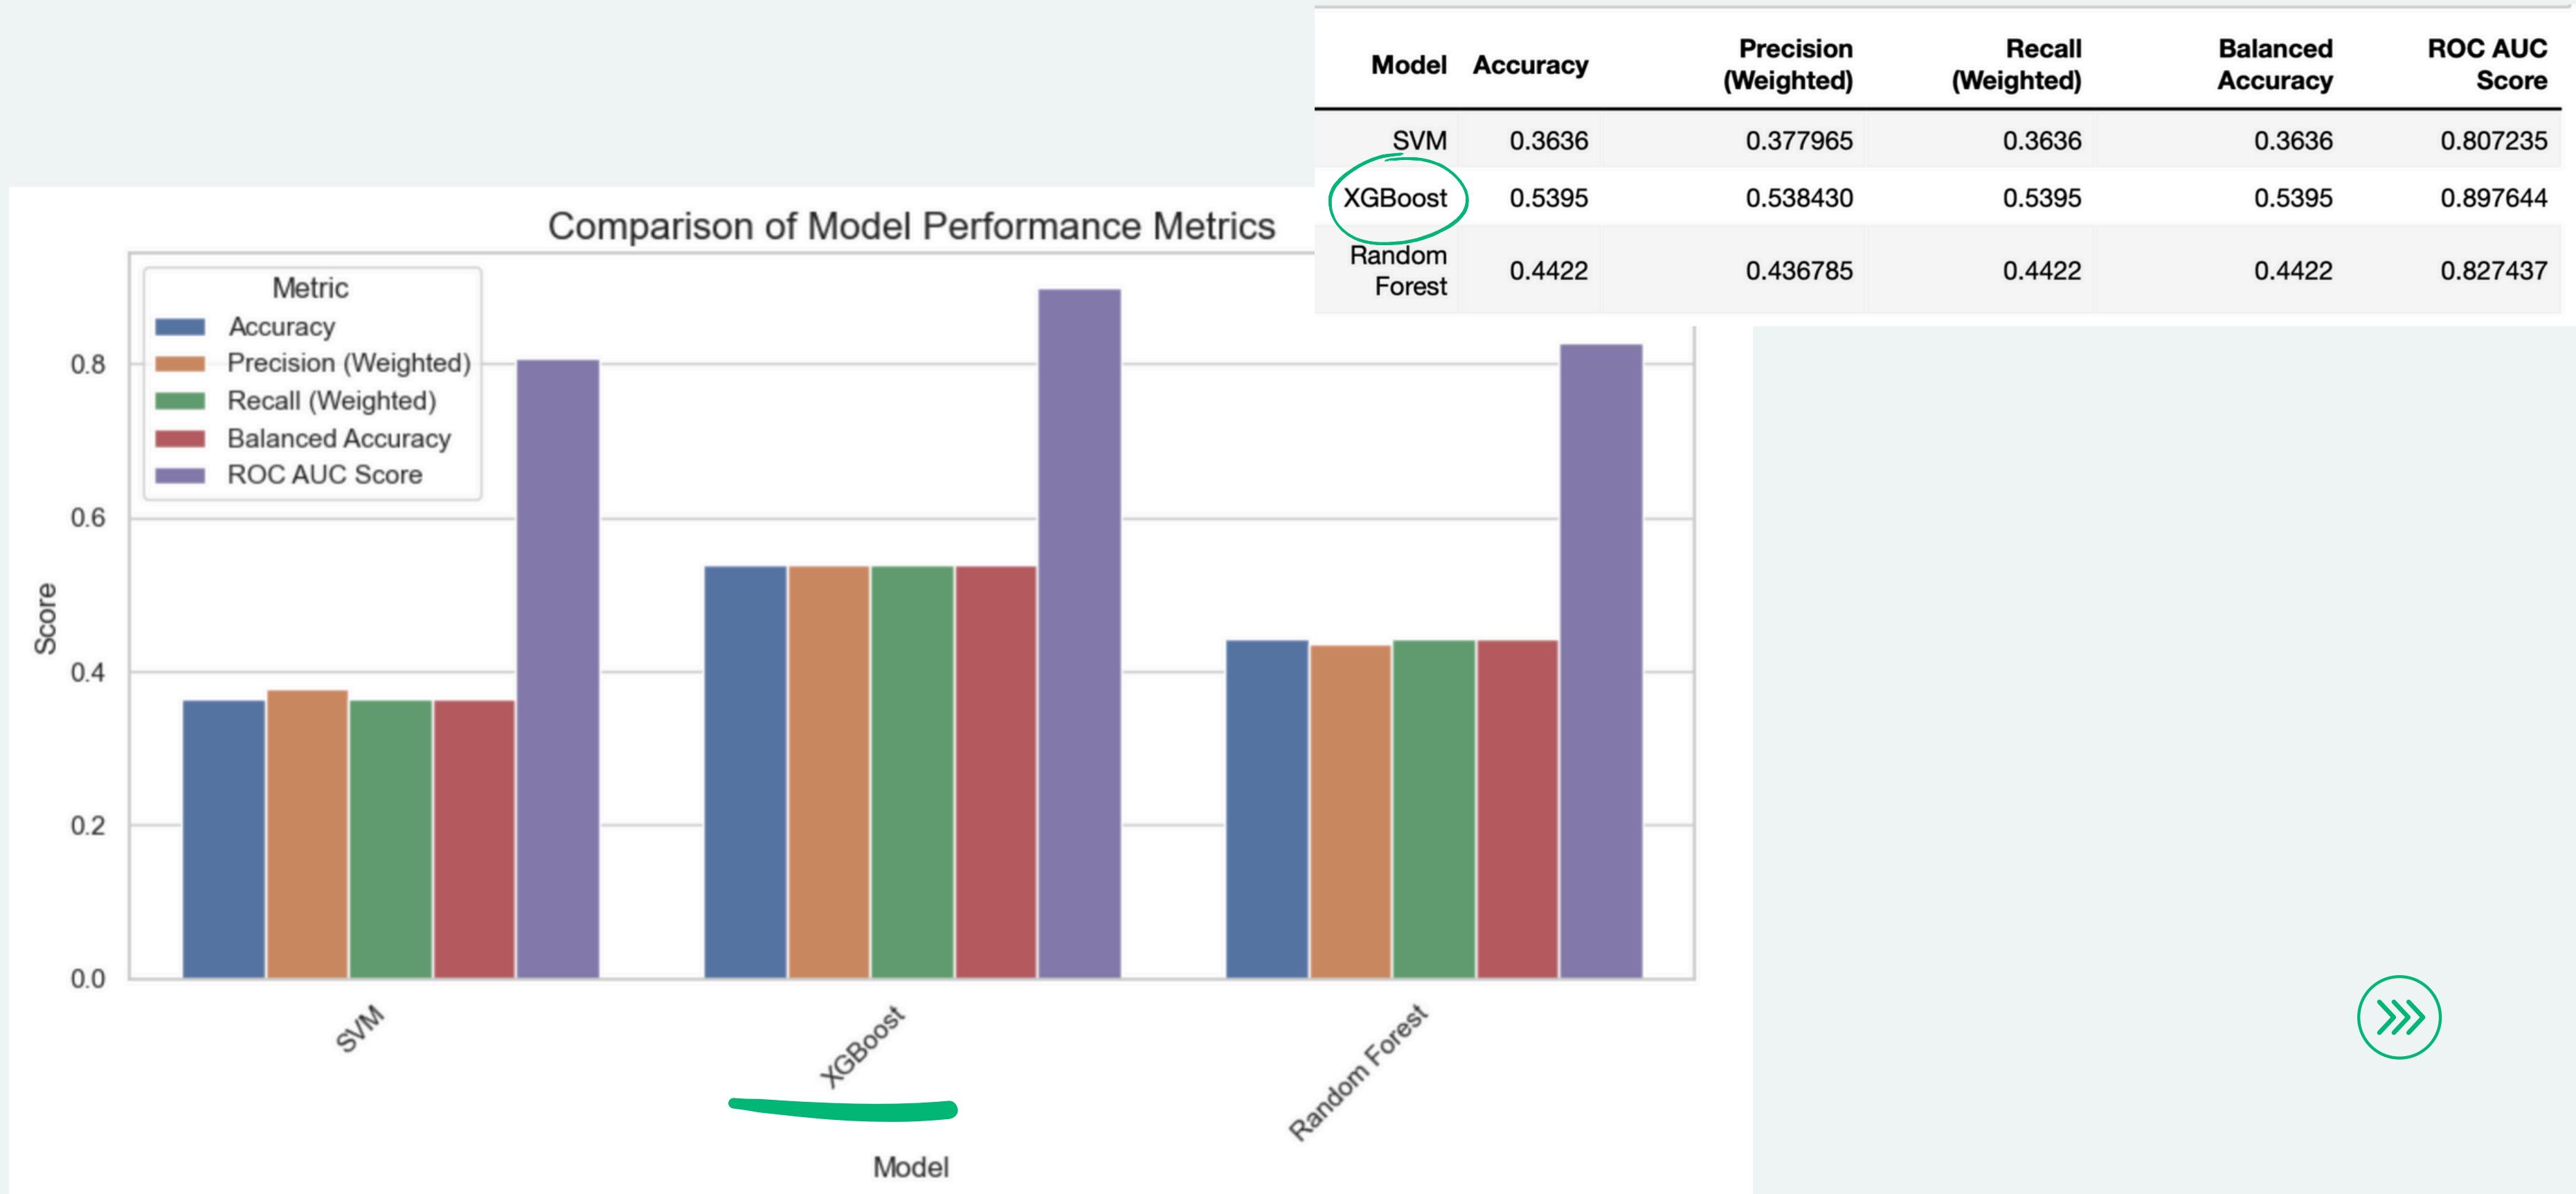
| Class      | Precision | Recall | F1-Score | Support |
|------------|-----------|--------|----------|---------|
| Airplane   | 0.64      | 0.63   | 0.63     | 1000    |
| Automobile | 0.66      | 0.68   | 0.67     | 1000    |
| Bird       | 0.49      | 0.41   | 0.45     | 1000    |
| Cat        | 0.41      | 0.35   | 0.38     | 1000    |
| Deer       | 0.47      | 0.52   | 0.49     | 1000    |
| Dog        | 0.46      | 0.46   | 0.46     | 1000    |
| Frog       | 0.56      | 0.68   | 0.61     | 1000    |
| Horse      | 0.63      | 0.57   | 0.6      | 1000    |
| Ship       | 0.61      | 0.62   | 0.62     | 1000    |
| Truck      | 0.65      | 0.68   | 0.66     | 1000    |



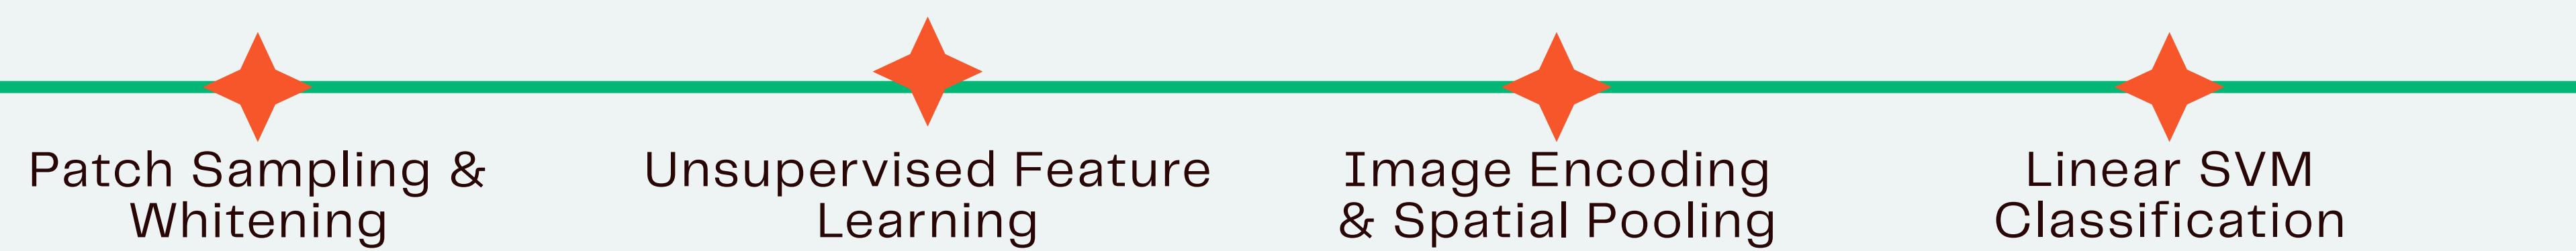
# COMPARISON OF THE MODELS



# COMPARISON OF THE MODELS



# CLASSIFICATION INTEGRATING K-MEANS CLUSTERING



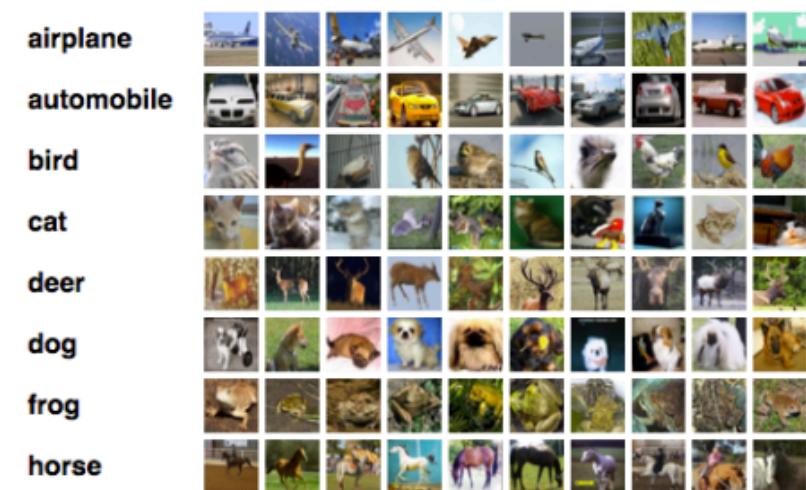
## Classifying CIFAR-10 Images Using Unsupervised Feature & Ensemble Learning

Truc Viet “Joe” Le  
Heinz College  
Carnegie Mellon University  
tjle@andrew.cmu.edu

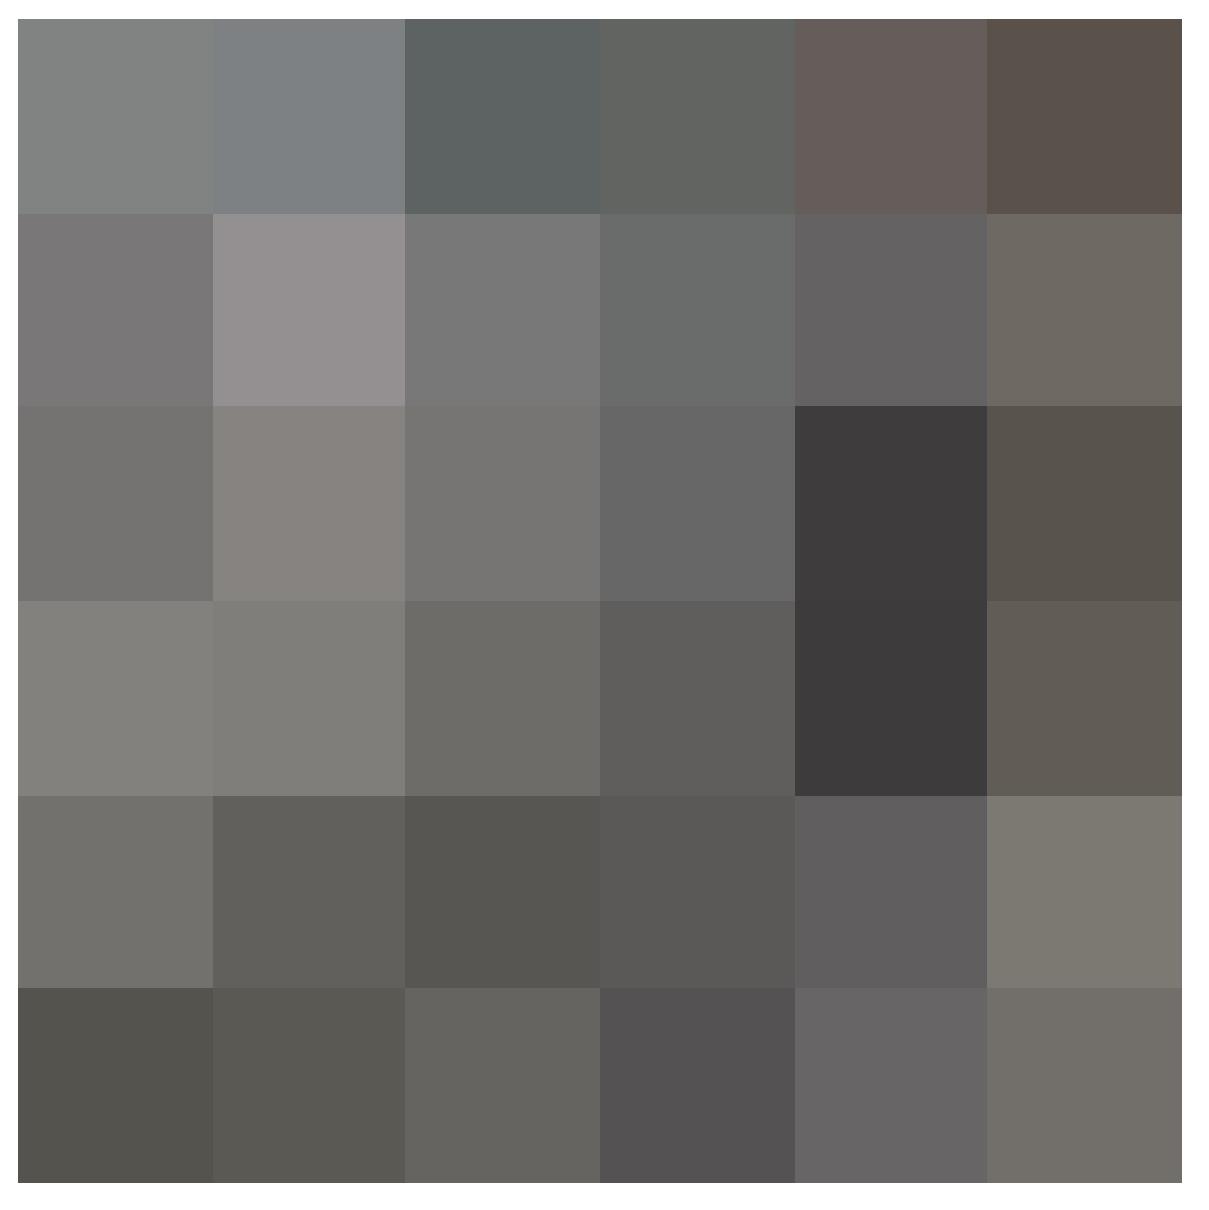
Naassih Gopee  
School of Computer Science  
Carnegie Mellon University  
ngopee@andrew.cmu.edu

### ABSTRACT

We perform the image classification task on the CIFAR-10 dataset, where each image belongs to one of the ten distinct classes. The classes are mutually exclusive and are mostly objects and animals. The images are small ( $32 \times 32$  pixels), of uniform size and shape, and RGB coloured. We implement an proposed image preprocessing framework to learn and extract the salient features of the images. The method was demonstrated to increase the classification performance significantly. We testify such claim and see a considerable improvement of more than 15% from the baseline (i.e., without preprocessing). We further experiment with various parameters and settings of the proposed method to tune the



# Patch Extraction



$(6, 6, 3)$

# Whitening + normalization

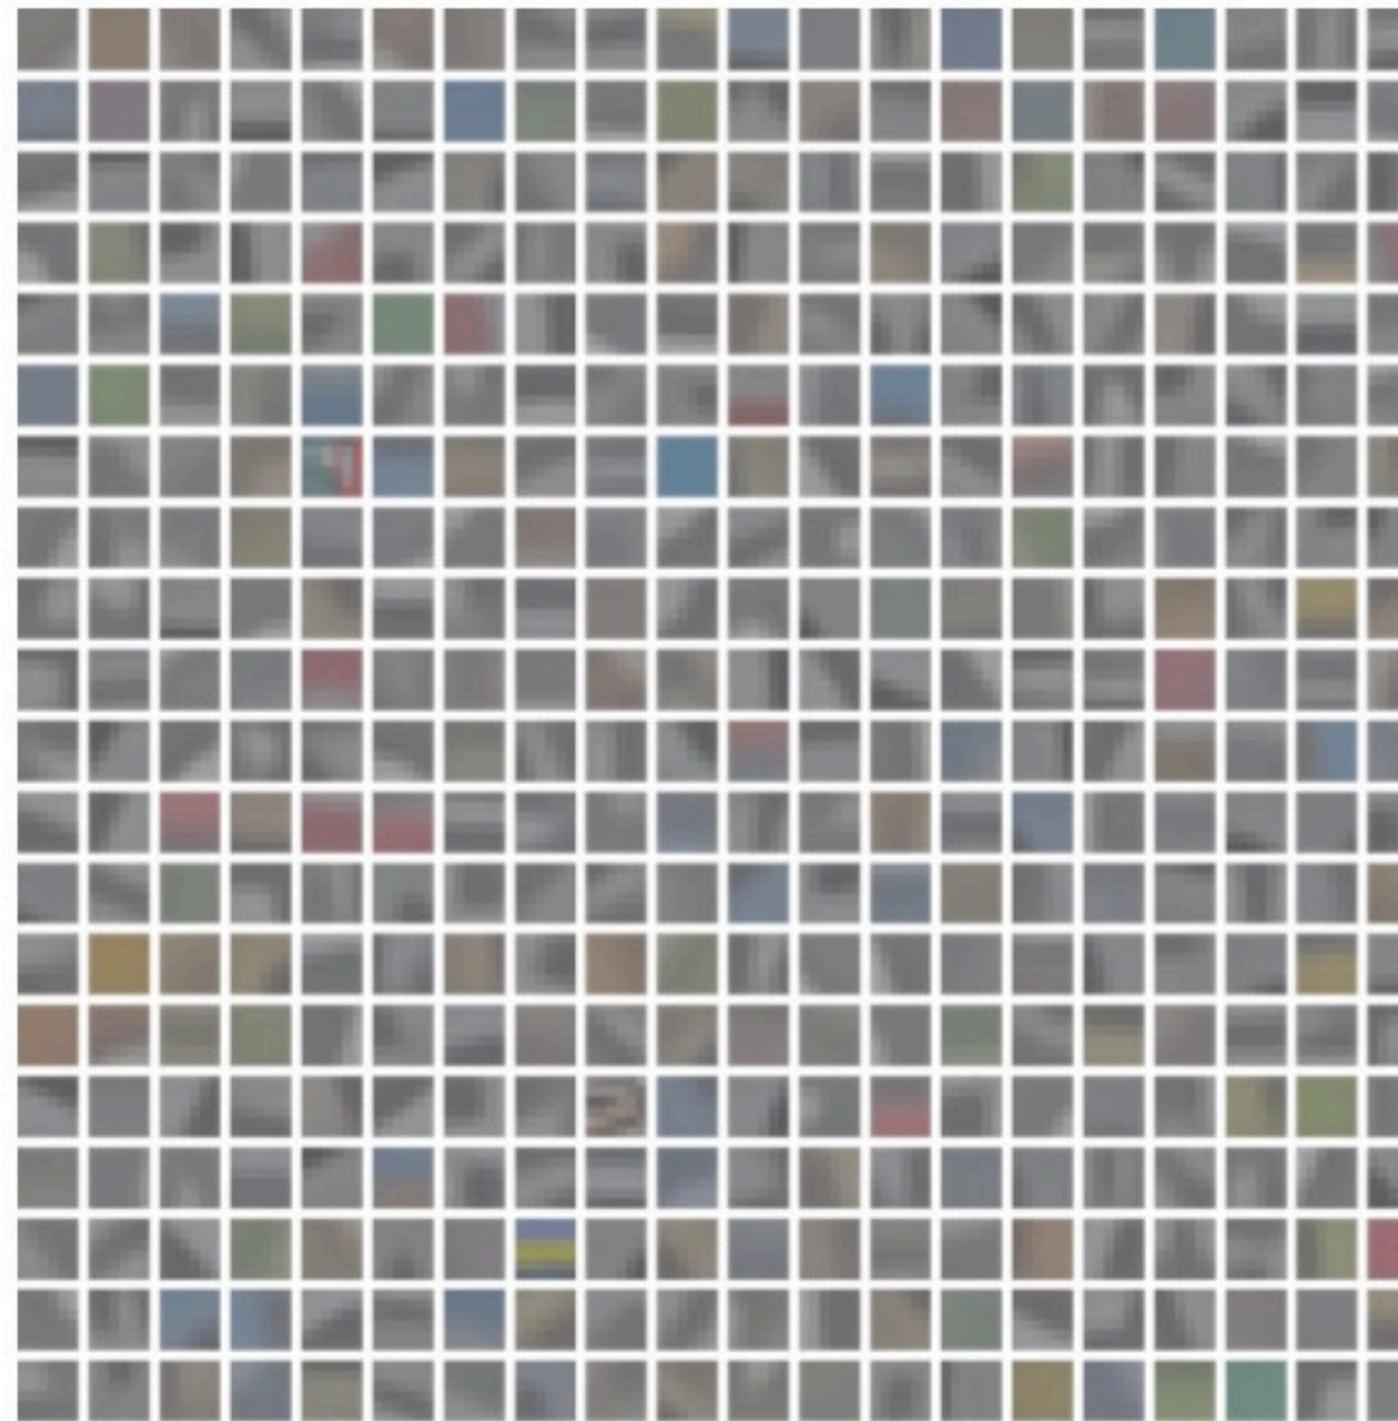


Image whitening – de-correlation transformation that transforms each patch into a matrix whose covariance matrix is the identity matrix.

Whitening amplifies variance along decorrelated directions – making patterns and contrasts more pronounced. That's why the whitened patch may appear more "colorful."

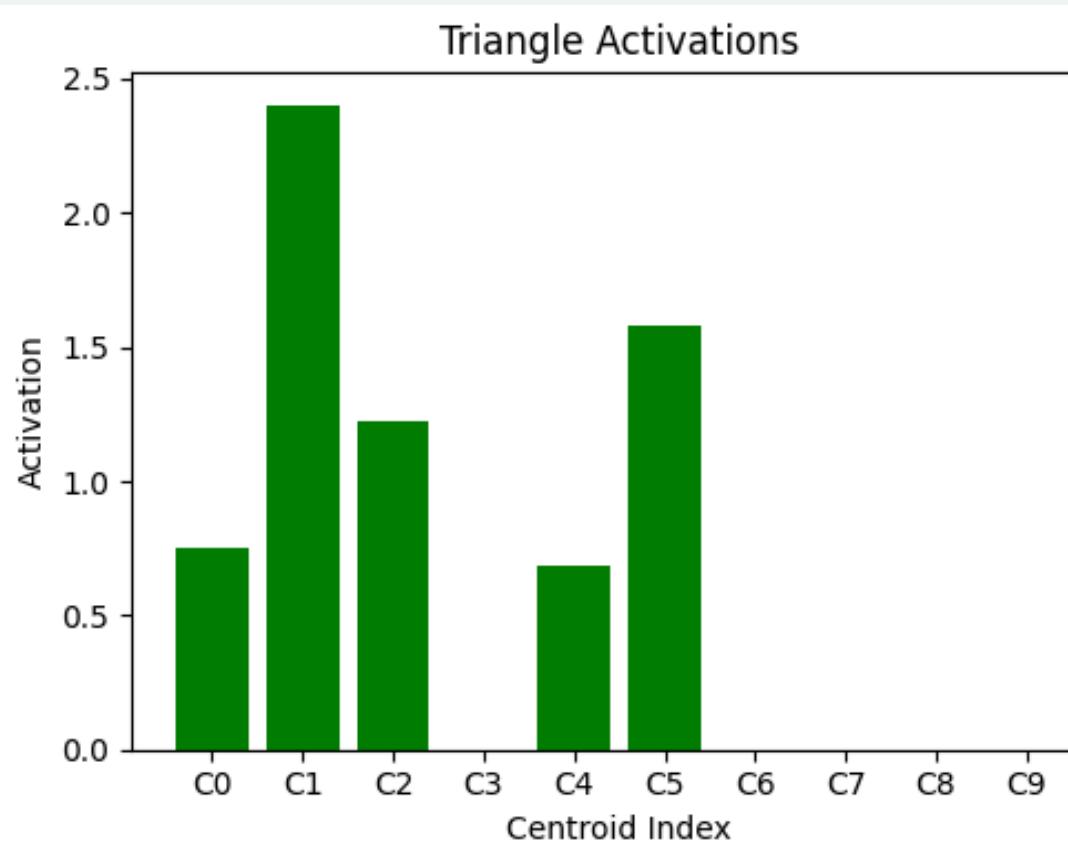
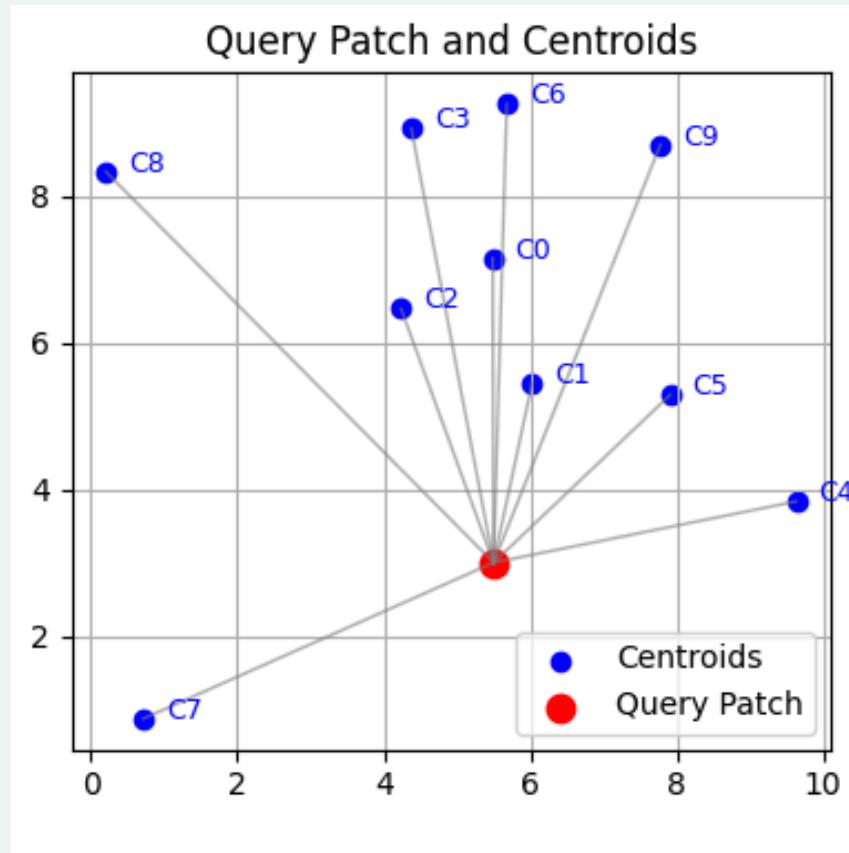
# K-means

Visualisation of the 400 centroids  
performed on 200000 random  
patches



# Triangle activation + pooling

Each patch is encoded using triangle activation with respect to the learned centroids. These patch responses are pooled across a  $3 \times 3$  grid, producing a final  $K \times 9$ -dimensional feature vector.



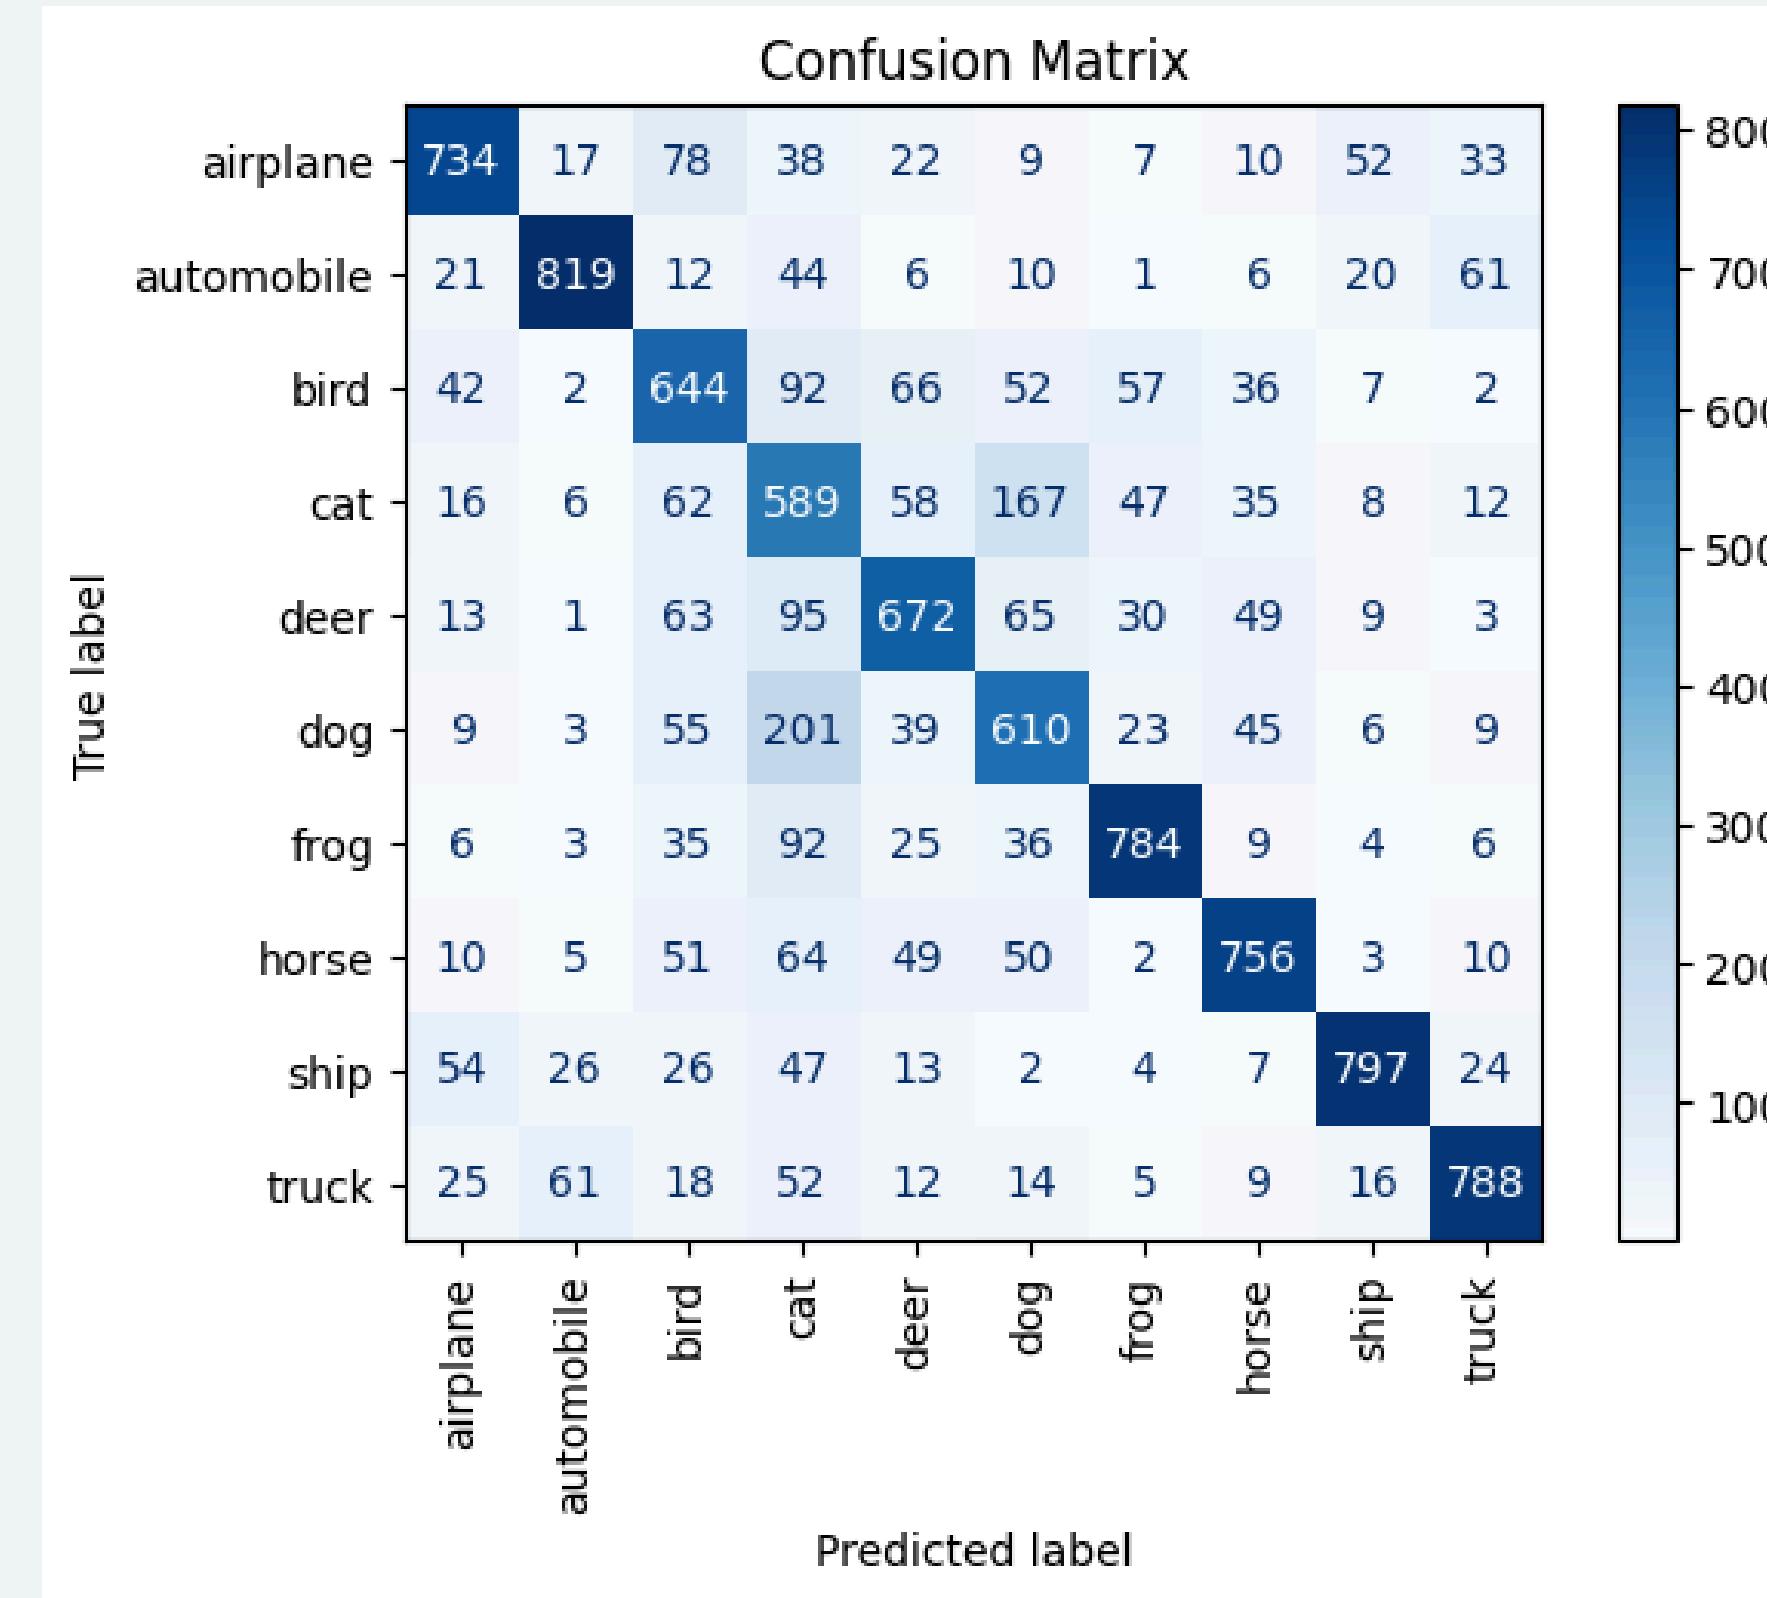
[0.0, 1.74, 0.0, 0.52, 0.0, ..., 0.0], shape for one image:  $(676, K)$



# Linear SVM

# K = 1600

Number of training samples: 50000  
Number of test samples: 10000



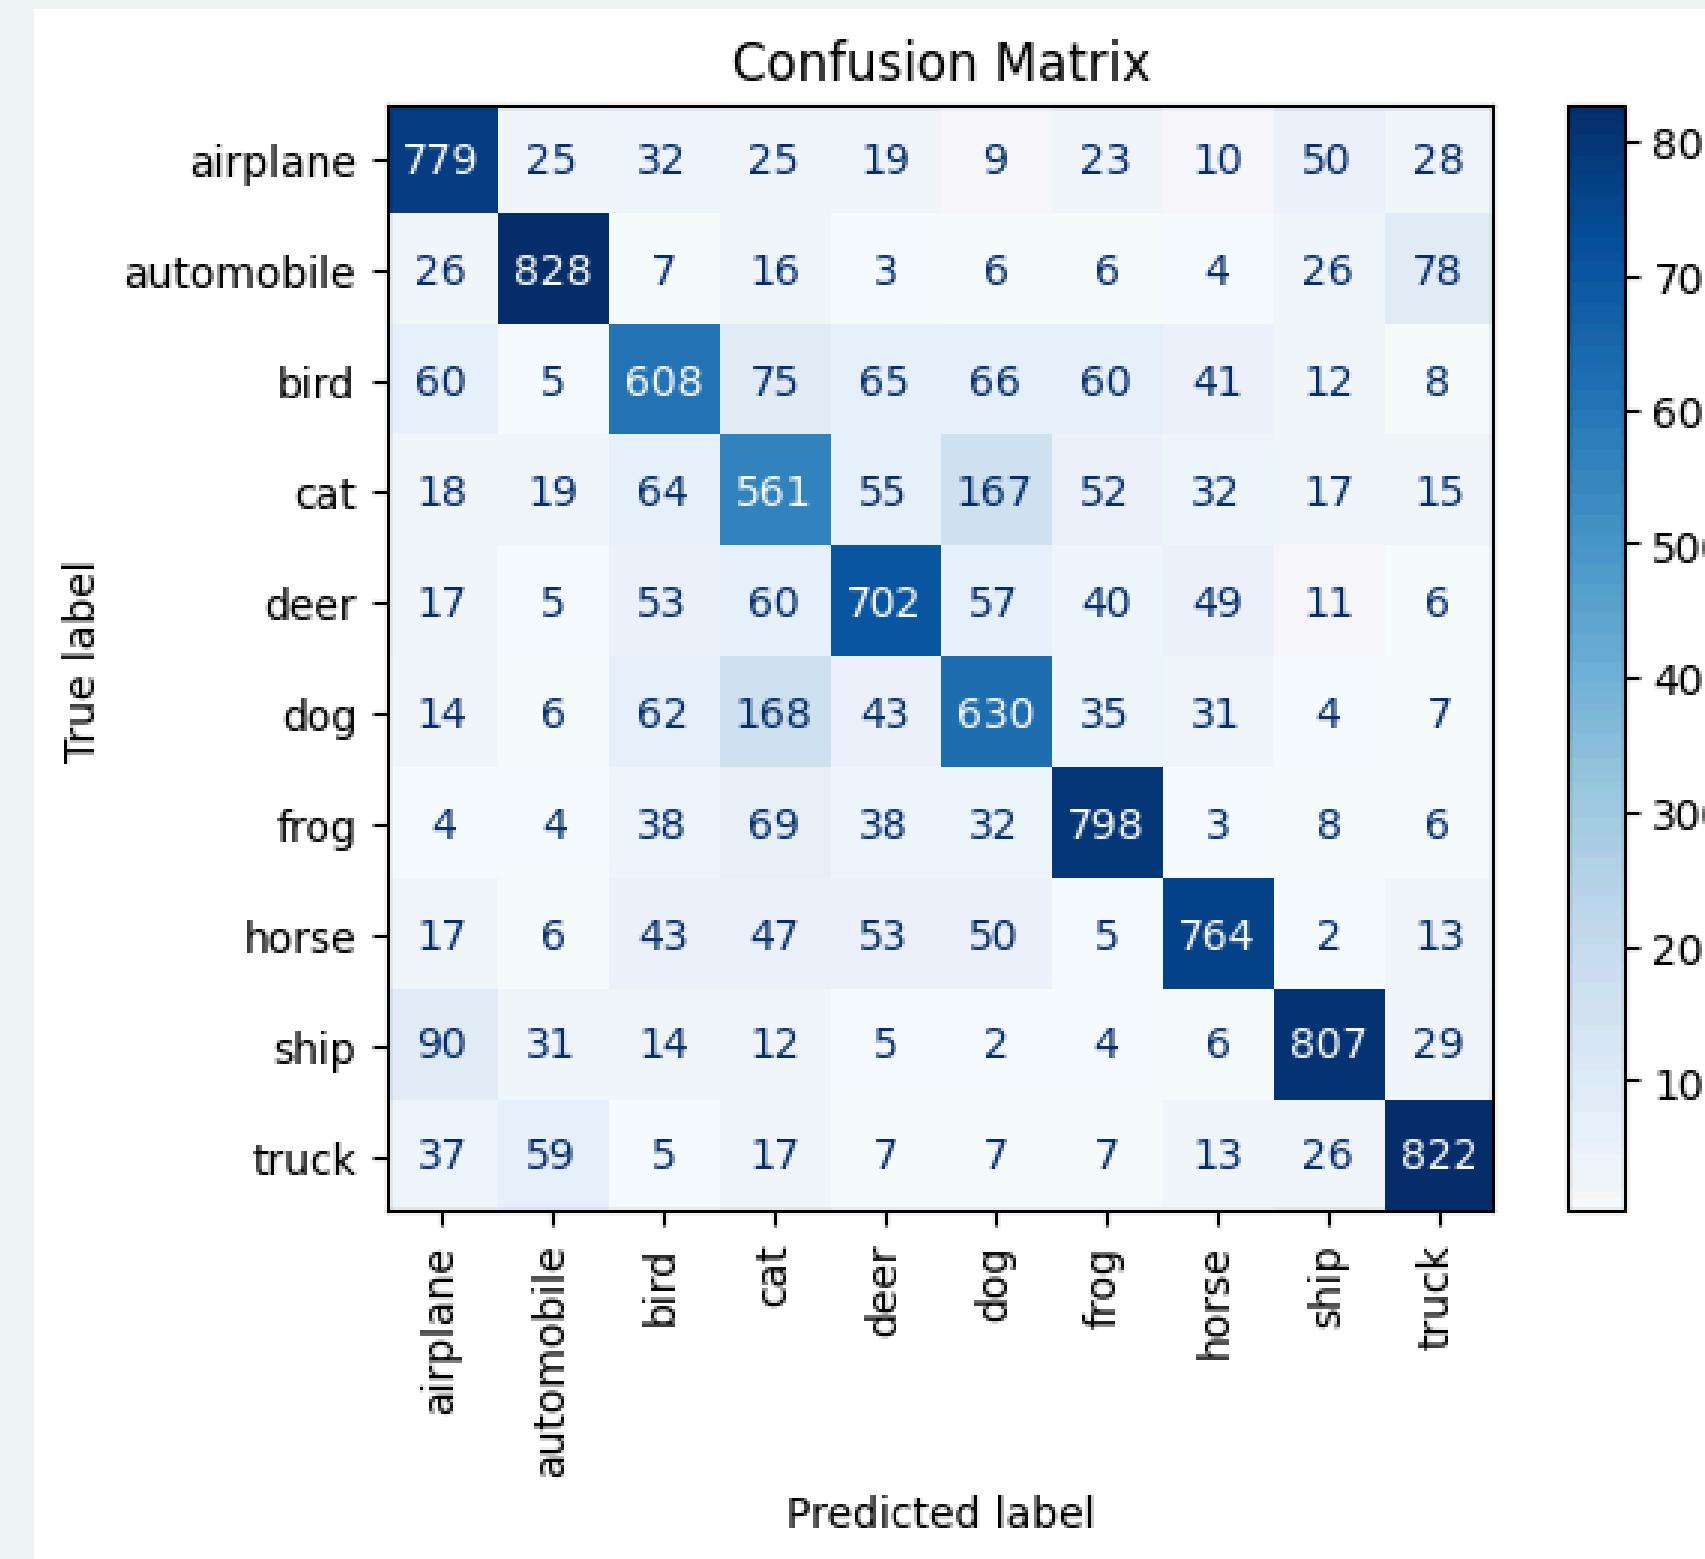
**Final Accuracy: 71.93%**

**automobile: 81.90%**  
**ship: 79.70%**  
**truck: 78.80%**  
**frog: 78.40%**  
**horse: 75.60%**  
**airplane : 73.40%**  
**deer: 67.20%**  
**bird: 64.40%**  
**dog: 61.00%**  
**cat: 58.90%**

# Linear SVM

**K = 4000**

Number of training samples: 35000  
Number of test samples: 10000

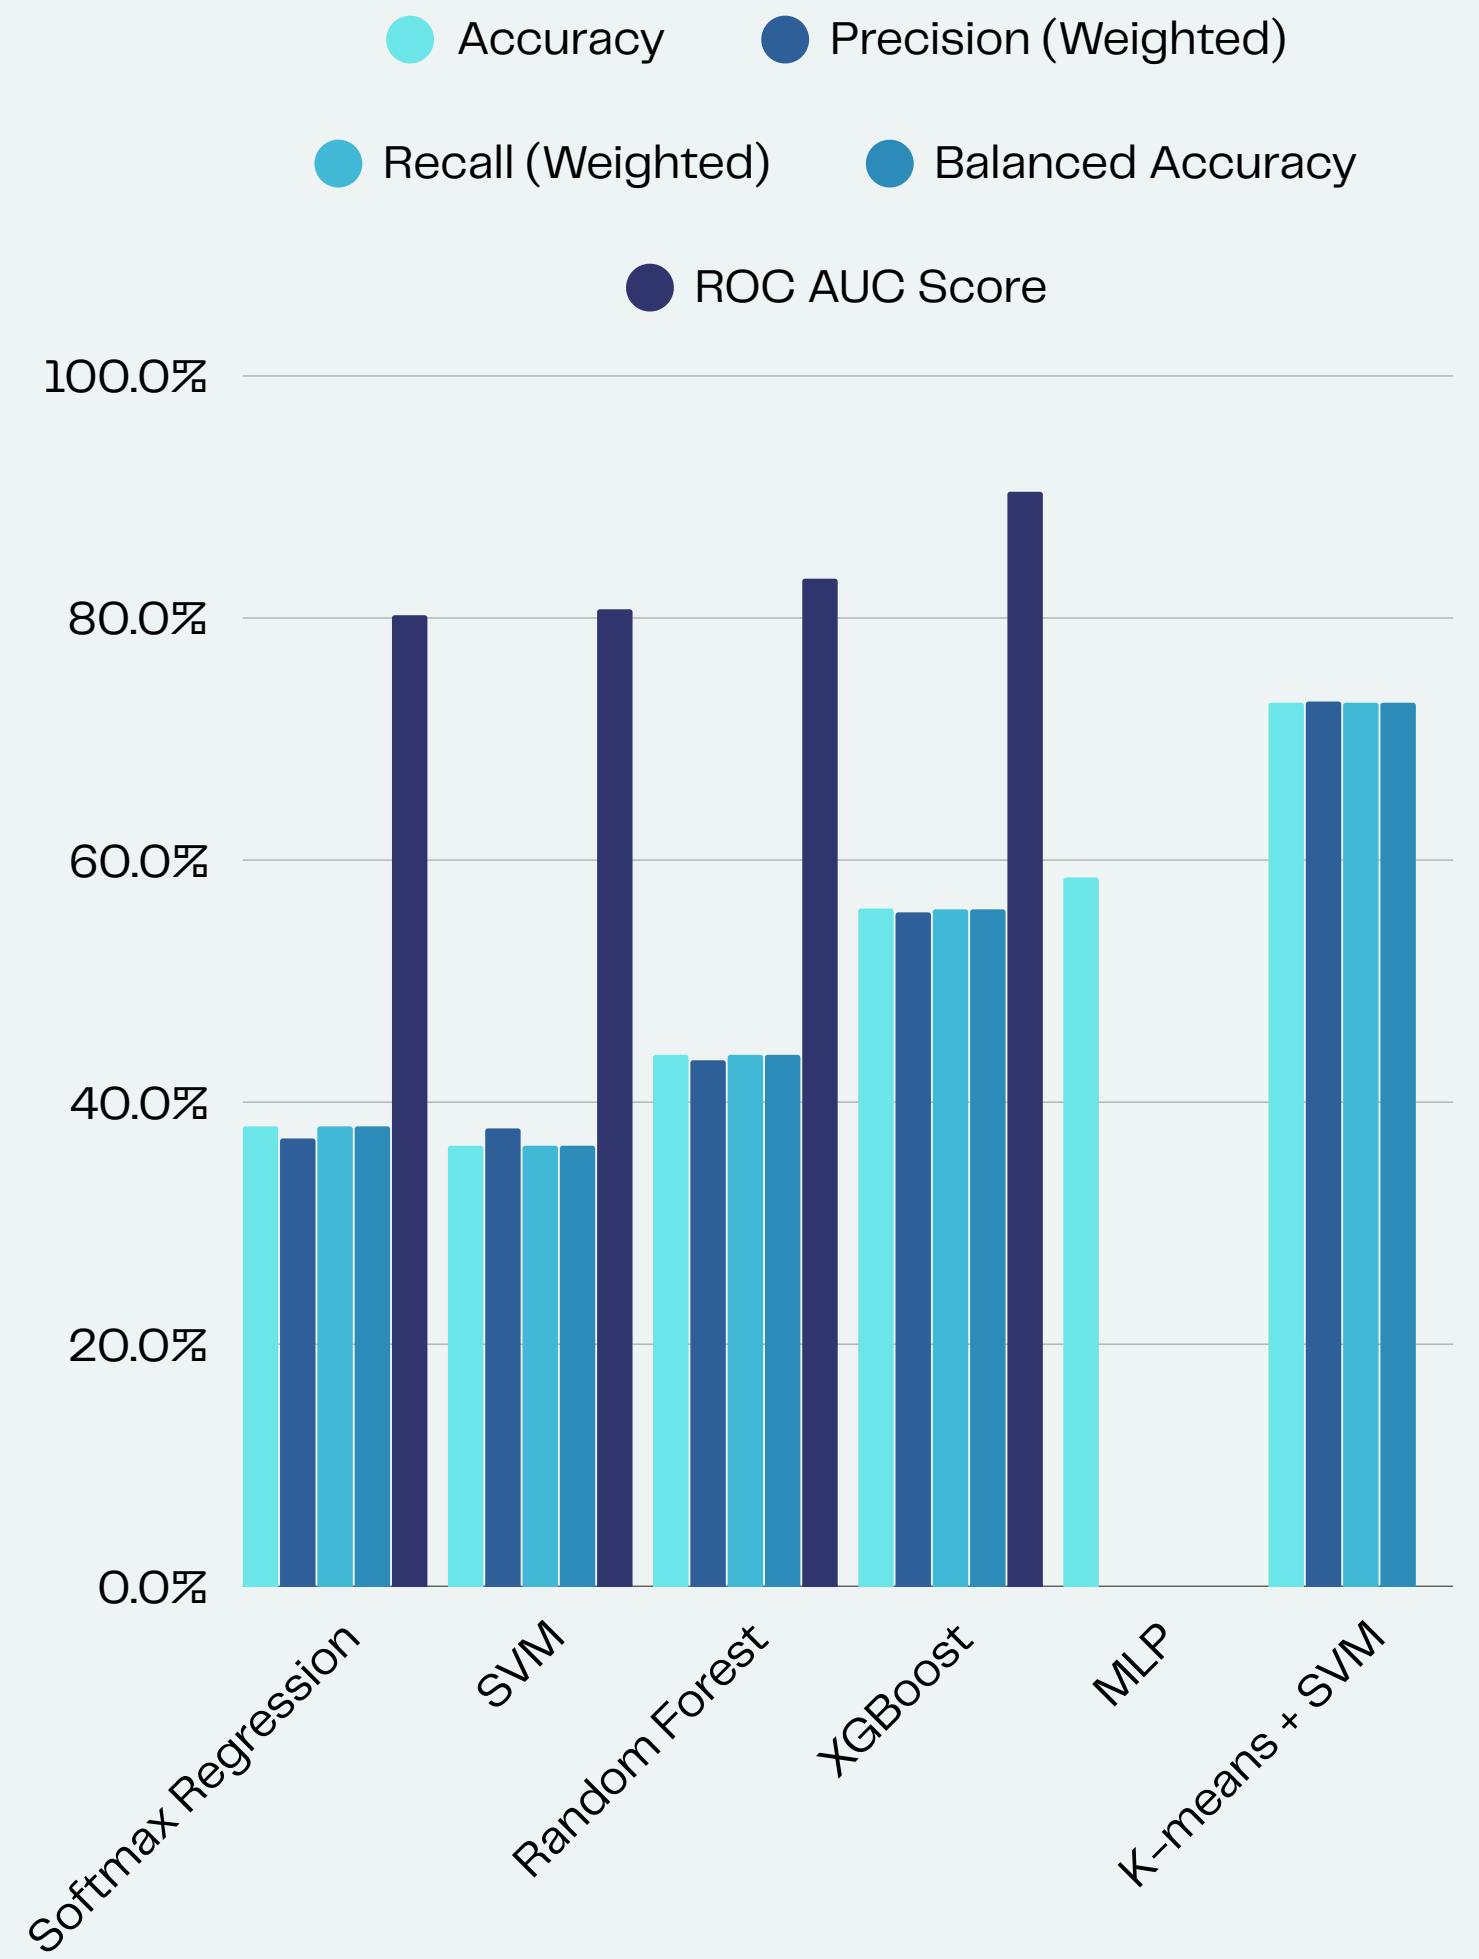


**Final Accuracy 72.99%**

**automobile: 82.80%**  
**truck: 82.20%**  
**ship: 80.70%**  
**frog: 79.80%**  
**airplane: 77.90%**  
**horse: 76.40%**  
**deer: 70.20%**  
**dog: 63.00%**  
**bird: 60.80%**  
**cat: 56.10%**

# OVERALL METRICS

| Model              | Accuracy | Precision (Weighted) | Recall (Weighted) | Balanced Accuracy | ROC AUC Score |
|--------------------|----------|----------------------|-------------------|-------------------|---------------|
| Softmax Regression | 38.00%   | 37.00%               | 38.00%            | 38.00%            | 80.23%        |
| SVM                | 36.40%   | 37.83%               | 36.40%            | 36.40%            | 80.72%        |
| Random Forest      | 43.91%   | 43.46%               | 43.91%            | 43.91%            | 83.25%        |
| XGBoost            | 56.00%   | 55.68%               | 55.93%            | 55.93%            | 90.43%        |
| MLP                | 58.56%   | —                    | —                 | —                 | —             |
| K-means + SVM      | 73.00%   | 73.10%               | 73.00%            | 73.00%            | —             |



# CONCLUSION



Based on the evaluation of multiple models, the **K-means + SVM** pipeline emerged as the most effective approach. While traditional classifiers like **Softmax Regression**, **SVM**, **Random Forest**, and **XGBoost** achieved moderate ROC AUC scores, their accuracy remained below satisfactory levels, with the best among them (XGBoost) reaching only 55.93% accuracy. The **MLP** model performed slightly better with 58.56% accuracy, but still fell short of expectations. In contrast, the K-means + SVM combination achieved a significantly higher accuracy of 73.00%, along with strong precision and recall scores, indicating consistent and reliable classification performance



# REFERENCES

**Le, T. V., & Gopiee, N. (n.d.). Classifying CIFAR-10 Images Using Unsupervised Feature & Ensemble Learning.** Carnegie Mellon University. Retrieved from <https://www.andrew.cmu.edu/user/tjle/>

**Coates, A., Lee, H., & Ng, A. Y. (2011). An analysis of single-layer networks in unsupervised feature learning.** In Proceedings of the Fourteenth International Conference on Artificial Intelligence and Statistics (AISTATS), 2011.

**Dalal, N., & Triggs, B. (2005). Histograms of Oriented Gradients for Human Detection.** In Proceedings of the IEEE Computer Society Conference on Computer Vision and Pattern Recognition (CVPR 2005), Vol. 1, pp. 886–893.

